

COMP/ELEC 429/556

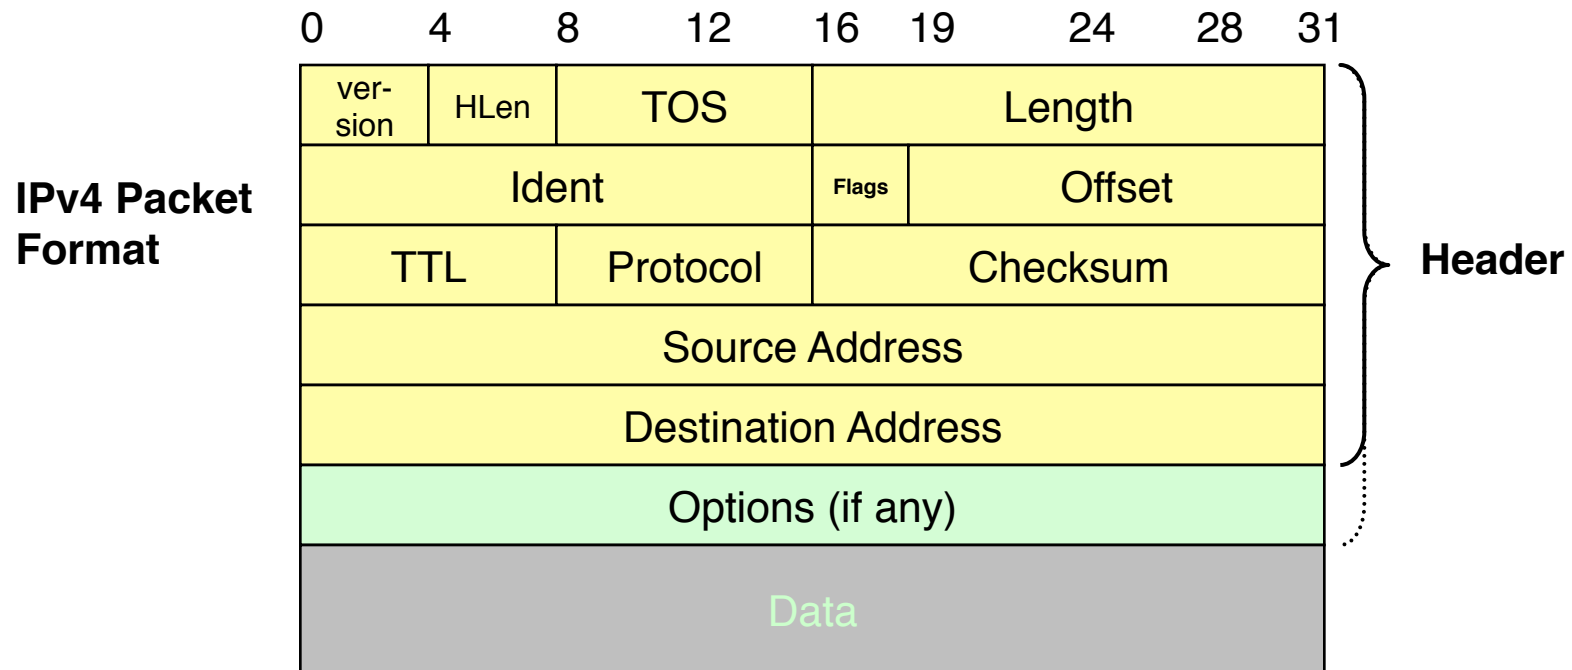
Introduction to Computer Networks

Domain Name System

Some slides used with permissions from Edward W.
Knightly, T. S. Eugene Ng, Ion Stoica, Hui Zhang

Data Packet

- Fundamental unit for communications in Internet



An IP address

2150268945

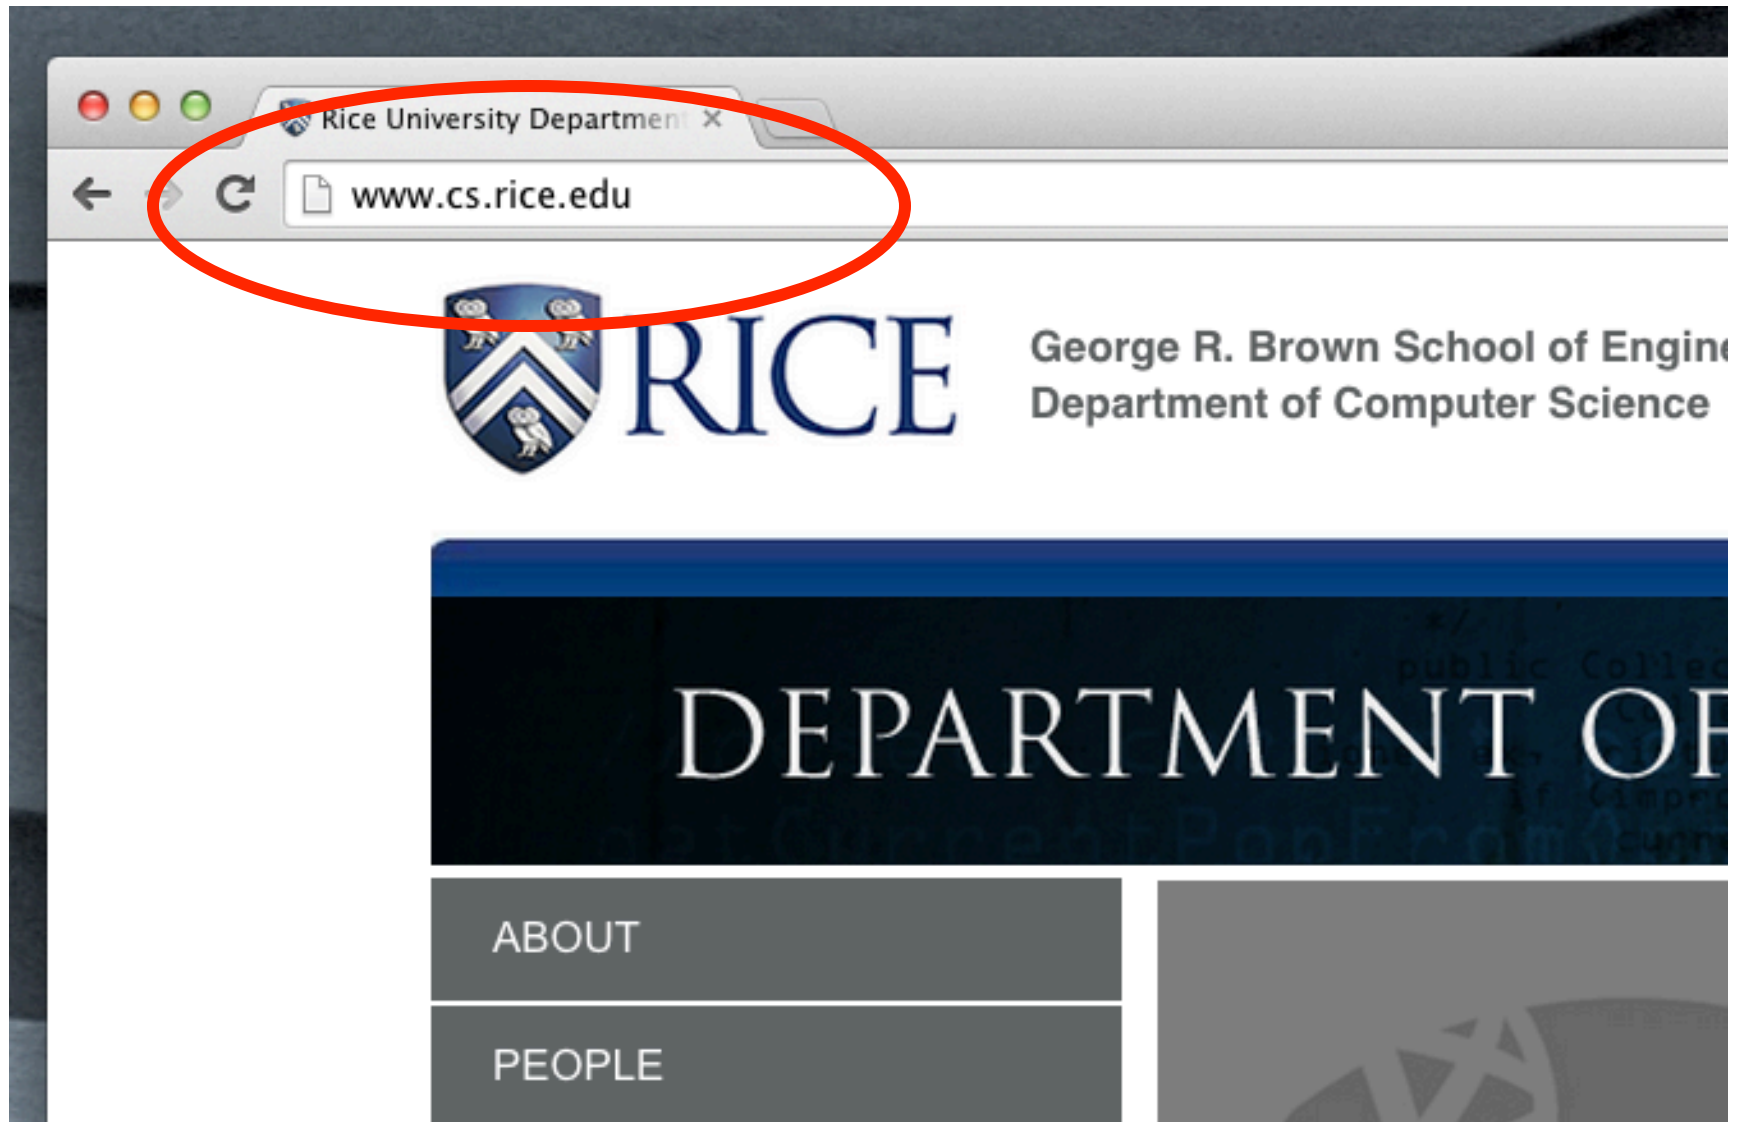
An IP address

128.42.128.17

10000000 00101010 10000000 00010001

A Domain Name

www.cs.rice.edu



Motivation

- Fact: A fundamental feature of the Internet is that every network interface is identified by a numerical IP address
- An application needs to know the IP address of the communication peer
- There is no magic, some out-of-band mechanism is needed
 - Word of mouth
 - Read it in magazine advertisements
 - Etc.
- But IP addresses are bad for humans to remember and tell each other, need names that makes sense to humans

Internet Names & Addresses

- Names: *e.g. www.rice.edu*
 - human-usable labels for machines
 - why this funny looking style with all these dots?
- Addresses: *e.g. 128.42.204.11*
 - 32-bit number, written this way for convenience
 - machine-usable labels for machines
 - more efficient to process than names
- How do you lookup from one to another?
- Let's try nslookup!

Remainder of this set of slides are for
your reference only. They will not be
needed for assignments.



Domain Name System is a case study of the importance of scalability




- Who's this guy?

Jon Postel

From Wikipedia, the free encyclopedia

Jonathan Bruce Postel ([/pəˈstɛl/](#); August 6, 1943 – October 16, 1998) was an American [computer scientist](#) who made many significant contributions to the development of the [Internet](#), particularly with respect to [standards](#). He is known principally for

Domain names used to be arbitrary and stored in one shared file hosts.txt



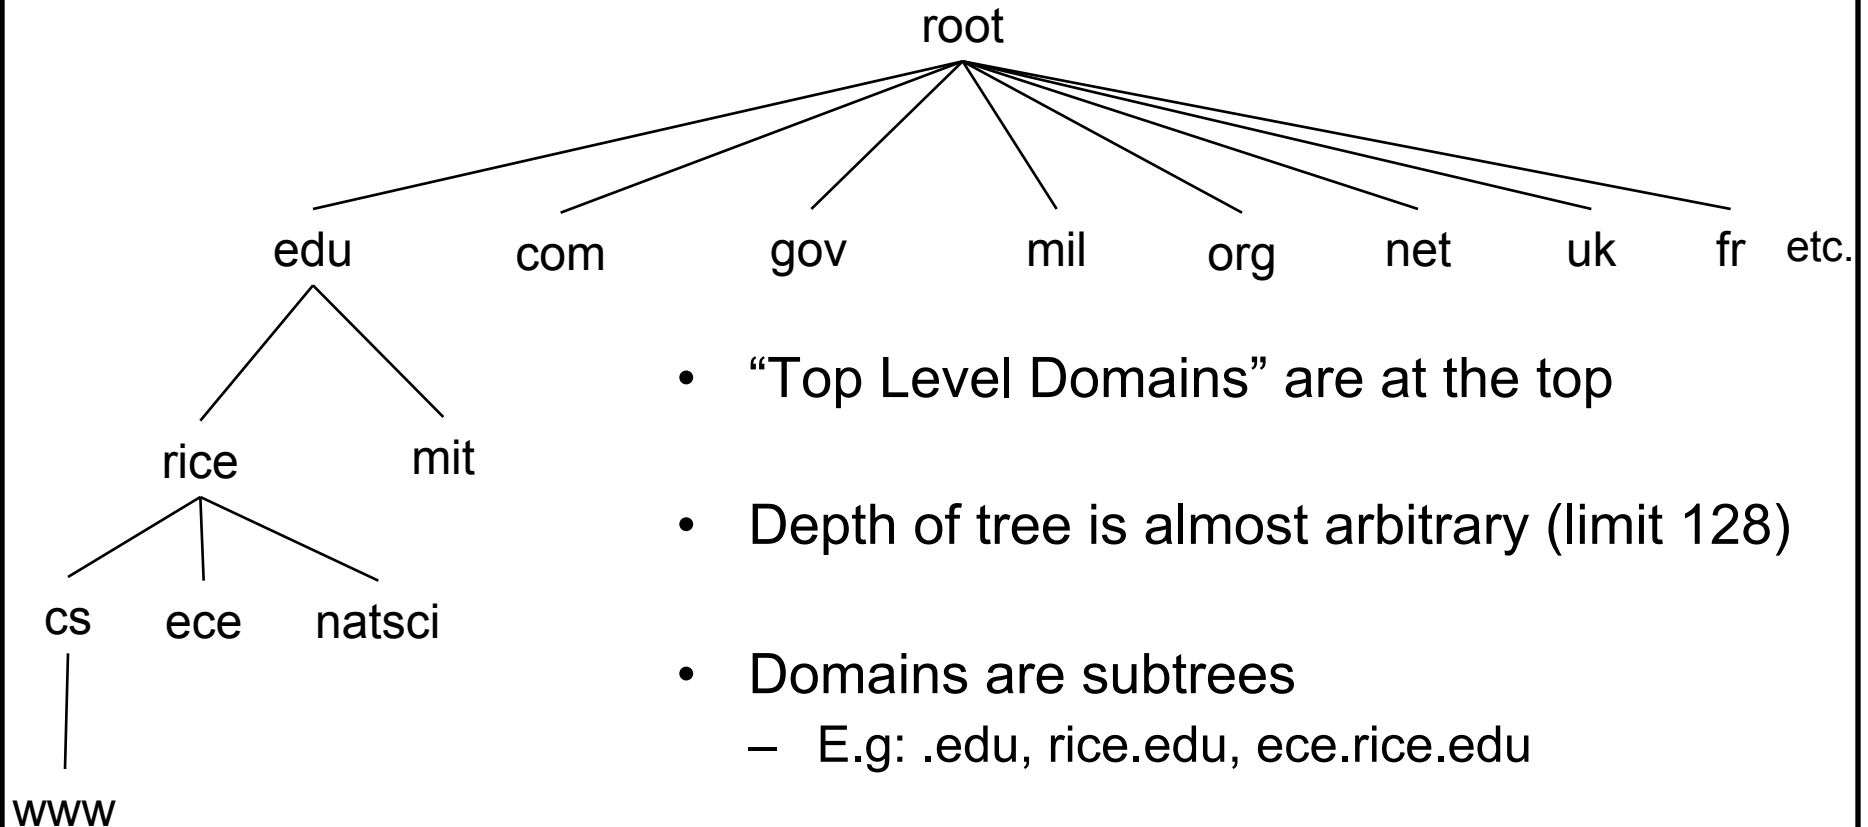
History

- Initially all host-address mappings were in a file called `hosts.txt` (in `/etc/hosts`)
 - Changes were submitted to SRI by email
 - New versions of `hosts.txt` ftp'd periodically from SRI
 - An administrator could pick names at their discretion
 - Any name is allowed: `eugenesdesktopatrice`
- As the Internet grew this system broke down because:
 - SRI couldn't handle the load
 - Hard to enforce uniqueness of names
 - Many hosts had inaccurate copies of `hosts.txt`
- How to build a lookup system that **scales!!!**
 - billions of names and addresses to insert/delete/modify
 - billions of lookups per second
- Strategy: Divide and Conquer!
 - Now do you see why domain names look like `www.rice.edu` instead of `lukeskywalker`?

Basic DNS Features

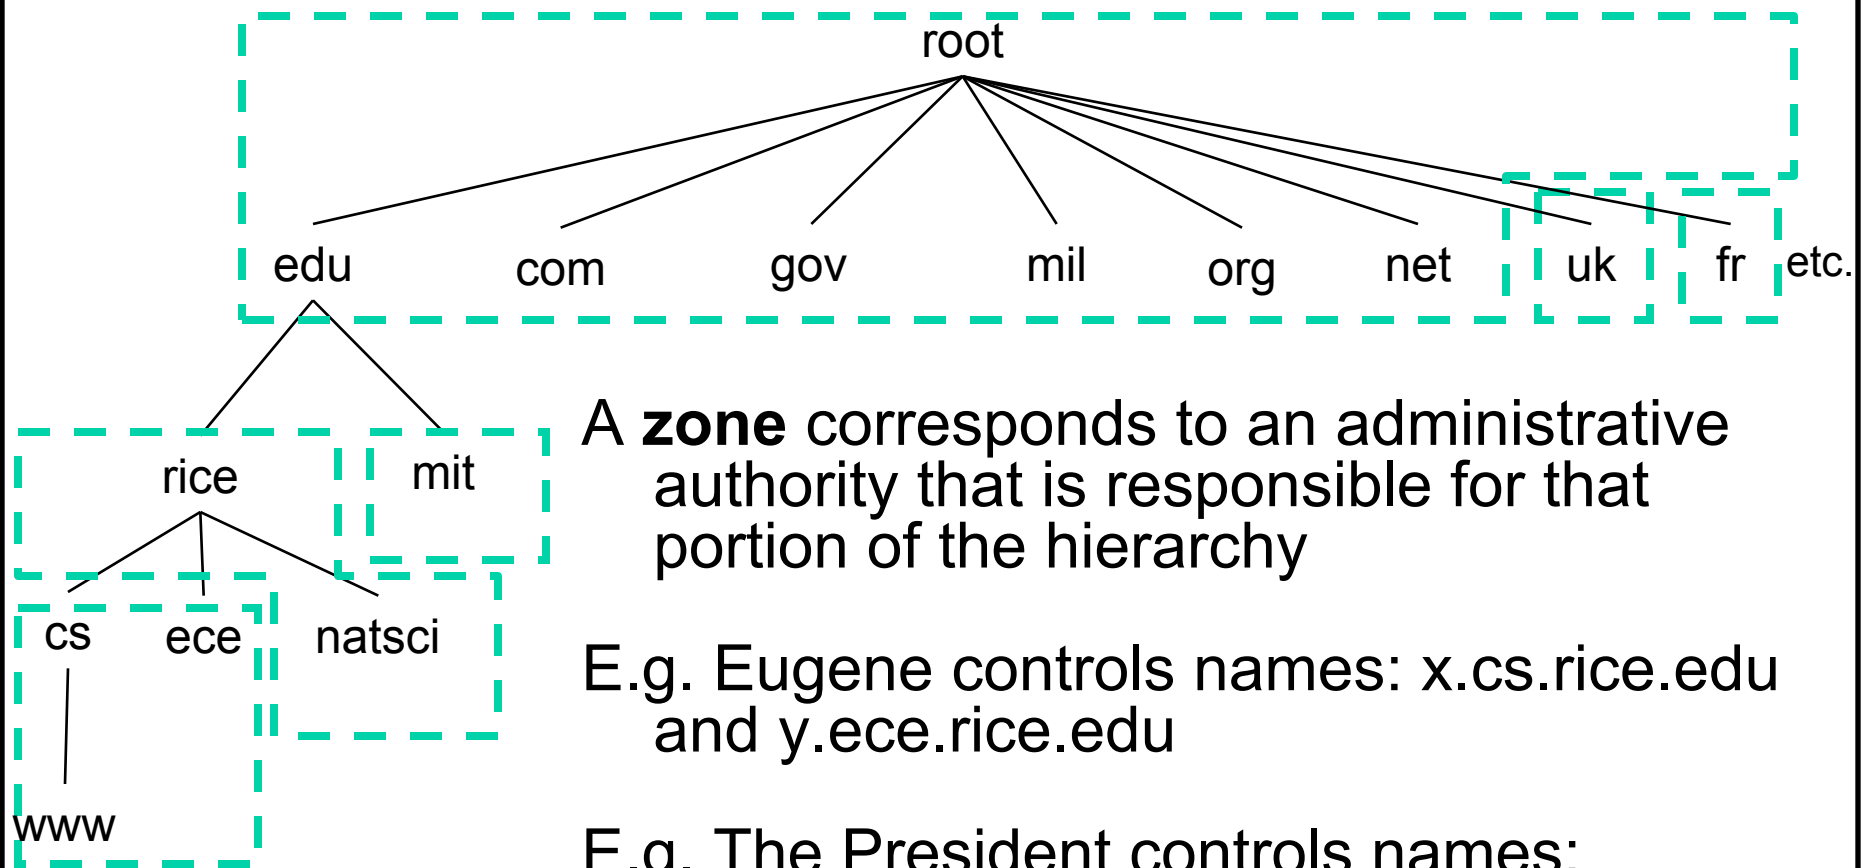
- Hierarchical namespace
 - as opposed to original flat namespace
- Distributed storage architecture
 - as opposed to centralized storage (plus replication)

Naming Hierarchy



- “Top Level Domains” are at the top
- Depth of tree is almost arbitrary (limit 128)
- Domains are subtrees
 - E.g: .edu, rice.edu, ece.rice.edu
- Name collisions avoided
 - E.g. rice.edu and rice.com can coexist, but uniqueness is job of domain

Host names are administered hierarchically



A **zone** corresponds to an administrative authority that is responsible for that portion of the hierarchy

E.g. Eugene controls names: x.cs.rice.edu and y.ece.rice.edu

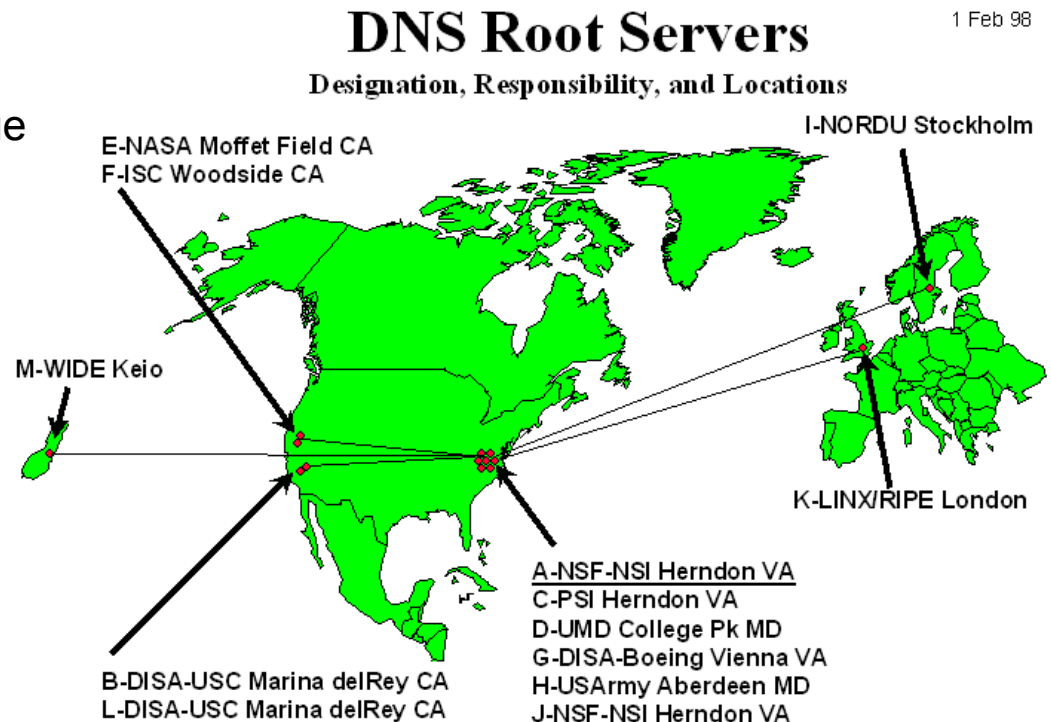
E.g. The President controls names: x.rice.edu and y.natsci.rice.edu

DNS Server Hierarchy

- Each server has authority over a portion of the hierarchy called zone
- Each server contains all the records for the hosts or domains in its zone
 - That zone can be empty; why will be revealed later
 - a server might be replicated for robustness
 - “Root server” knows about all top-level domains


DNS: Root Servers

- About a dozen root server IP addresses
 - Each address can refer to a large cluster of replicated servers to achieve sufficient performance
- Contacted by other servers that cannot resolve name



Load Balancers Are Very Common

- e.g. Amazon AWS



Basic Domain Name Resolution

- Every host knows a local DNS server
 - Through DHCP, for example
 - Sends all queries to a local DNS server
- Every local DNS server knows the ROOT servers
 - When no locally cached information exists about the query, talk to a root server, and go down the name hierarchy from the root
 - If we lookup `www.rice.edu`, and we have a cached entry for the `rice.edu` name server, then we can go directly to the `rice.edu` name server and bypass the root server


Example of Iterated DNS Query

Root name server:

- May not know authoritative name server
- May know **intermediate name server**: who to contact to find authoritative name server

Iterated query:

- Contacted server replies with name/address of another server
- “I don’t know this name, but ask this server”




Example of Recursive DNS Query

Which scales better?
Iterative or Recursive?

Recursive query:

- Puts burden of name resolution on contacted name server
- Extra “**state**” maintained at the intermediate servers
- More “**state**” is more complexity and hurts scalability
- Support is uncommon



DNS Resource Records

- DNS Query:
 - Two fields: (name, type)
- Resource record is the response to a query
 - Four fields: (name, value, type, TTL)
 - TTL (time to live) is the duration the response can be cached
 - There can be multiple valid responses to a query
- Type = A:
 - name = hostname
 - value = IP address

Other Common Record Types

- Type = NS:
 - name = domain
 - value = name of dns server for domain
- Type = CNAME:
 - name = hostname
 - value = canonical name
- Type = MX:
 - name = domain in email address
 - value = canonical name of mail server and priority

Discussions

- DNS caching
 - Each record has a TTL
 - Crucial to scalability
 - Improve performance by saving results of previous lookups
 - E.g. results of address records and name server records (e.g. if rice.edu name server is cached, then can bypass root server the second time looking for a rice.edu host)
- DNS “hacks”
 - Return records based on requesting IP address
 - Round-robin over a list of IP addresses mapped to the same name for load balancing
 - Return address of least loaded machine
 - Basis for many web content distribution networks such as Akamai and Limelight

Discussions

- DNS has been a large source of security problems throughout the history of the Internet

Examples:

- DNS servers can be leveraged to amplify a denial of service attack
- DNS cache poisoning can allow attackers to direct traffic to malicious machines
- Original DNS server software had lots of buffer overflow bugs that could be exploited by an attacker to take over the DNS server