

# COMP/ELEC 429/556

## Introduction to Computer Networks

Creating a Network Application


Some slides used with permissions from Edward W.  
Knightly, T. S. Eugene Ng, Ion Stoica, Hui Zhang

# How to Programmatically Use the DNS System?

Operating system comes to the rescue

- `gethostbyname()`
- `gethostbyaddr()`
- These functions are not re-entrant
  - Cannot have 2 program threads calling them concurrently
  - Call them in your main thread or use the more complicated thread-safe alternatives `gethostbyname_r()` and `gethostbyaddr_r()`
- On CLEAR, you can type “man `gethostbyname`” to get a detailed description of the function
  - “man” stands for manual

# How to Programmatically Send/Receive Data over the Internet?




# Operating System comes to the rescue

- Reliable byte stream network communications service
  - The most common model
  - Underlies almost every network application you use

# Reliable Byte Stream Model


- The elemental transmission unit is a byte (8 bits)
- Bytes are delivered reliably, first in first out
  - We will learn how this is accomplished later




# Reliable Byte Stream

- Bytes are delivered without any notion of application message units

Example 1



Example 2




## Application decides message format and how to interpret the byte stream

E.g.

- Each message is a null (0x00) terminated string
- In this case, a program reads in bytes until encountering 0x00 to form a valid message

# Project 1 Ping/Pong Message



2 size bytes  
0xf2 0x12



What is the size?

0x12f2 = 4,850

0xf212 = 61,970

Network byte order (a.k.a. Big endian) – first byte sent/received is the most significant

size is 0x12f2 = 4,850



# Why the need to specify network byte ordering convention for numbers?

- CPUs may choose little-endian or big-endian format when storing numbers in registers and RAM
  - Most Intel CPUs are “little-endian” (1<sup>st</sup> byte least significant)
  - Network byte ordering convention is “big-endian” (1<sup>st</sup> byte most significant)
- The number 4,850 (0x12f2), which is stored in “big-endian” ?



## What about bit ordering?

- Which bit in a byte received over the network is most significant?
- OMG, not again!




- Bit ordering within a byte is handled entirely in hardware

# Byte ordering conversion functions

- host order to network order
  - htons() for short integer (i.e. 16 bit)
  - htonl() for long integer (i.e. 32 bit)
  - htobe64() for 64 bit integer
- network order to host order
  - ntohs() for short integer (i.e. 16 bit)
  - ntohl() for long integer (i.e. 32 bit)
  - be64toh() for 64 bit integer
- Must be careful in deciding whether conversion is required
  - Is a variable numeric?
  - Is a numeric variable already stored in network byte order?
- These functions are “no-op” on CPUs that use network byte ordering (big-endian) for internal number representation
  - But a program must compile and work regardless of CPU, so these functions have to be used for compatibility

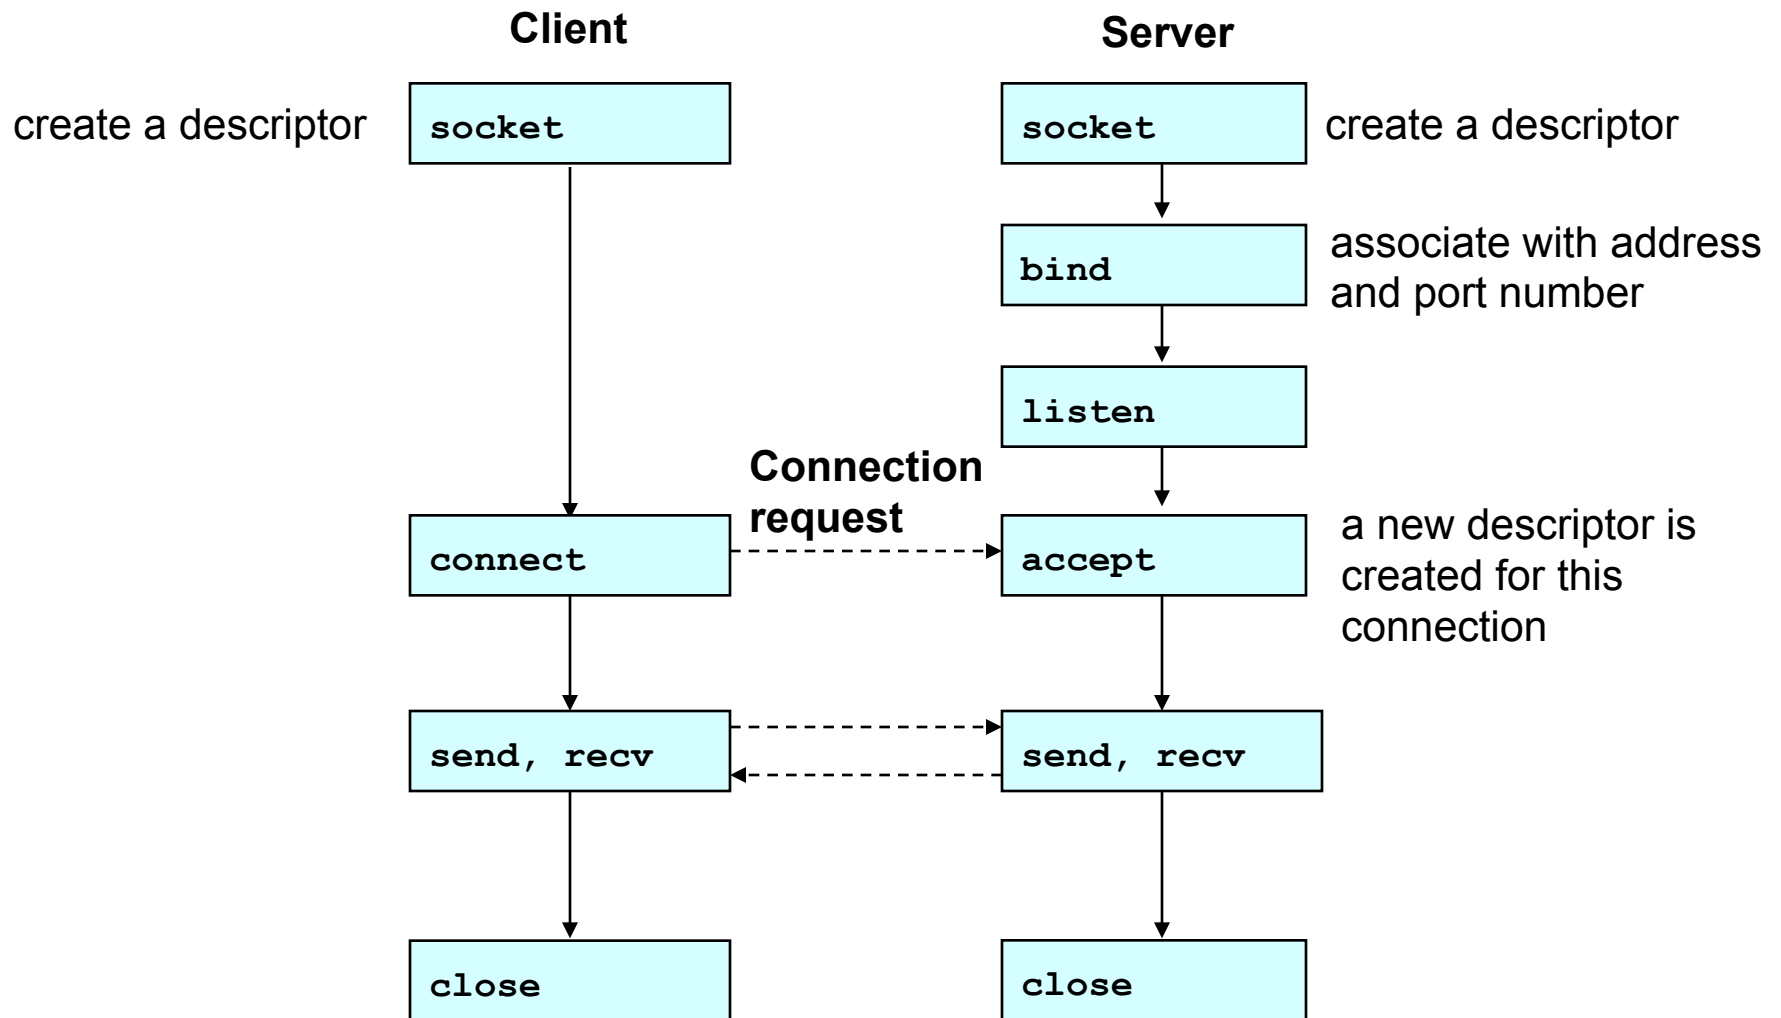
# How to Programmatically Send/Receive Data over the Internet?




# Operating System comes to the rescue

- Socket API
  - Refresh your memory of the mechanics of using socket for network communications

# Overview of the Socket API



# How to support multiple simultaneous connections?



# Operating System comes to the rescue

- Event-driven concurrency
  - Refresh your memory of one way to create the illusion of concurrently handling multiple network conversations even on a uni-processor
- Another way is multi-threading
  - Can achieve real concurrency on multi-processors
  - Refer to COMP321 material
- A combination of event-driven concurrency and multi-thread concurrency is needed to achieve the highest server scalability





# Event-Based Concurrent Servers

- Maintain a pool of descriptors
- Repeat the following forever:
  - Block until:
    - New connection request arrives on the listening descriptor
    - New data arrives on an existing connected descriptor
    - A connected descriptor is ready to be written to
    - A timeout occurs (can be configured to have no timeout)
  - If new connection request, add the new connection to the pool of connections
  - If new data arrives, read any available data from the connection
  - If descriptor is ready to be written, write whatever data is pending
  - If timeout, do whatever is appropriate for timeout
- Can wait for input from local I/O (standard input) and remote I/O (socket) simultaneously!
- Many implementations: `select()`, `poll()`, `epoll()`, etc.

# Event-Based Concurrent I/O

```
int select(int nfd, fd_set *readfds, fd_set *writefds,  
          fd_set *exceptfds, struct timeval *timeout);
```

- `readfds, writefds`
  - Opaque bit vector (max `FD_SETSIZE` bits) that indicates membership in a descriptor set
  - If bit  $k$  is 1, then descriptor  $k$  is a member of the descriptor set
- `nfd`
  - Maximum descriptor value + 1 in the set
  - Tests descriptors 0, 1, 2, ..., `nfd` - 1 for set membership
- `select()` returns the number of ready descriptors and sets each bit of `readfds, writefds` to indicate the ready status of its corresponding descriptor

## Macros for Manipulating Set Descriptors

- `void FD_ZERO(fd_set *fdset);`
  - Turn off all bits in `fdset`
- `void FD_SET(int fd, fd_set *fdset);`
  - Turn on bit `fd` in `fdset`
- `void FD_CLR(int fd, fd_set *fdset);`
  - Turn off bit `fd` in `fdset`
- `int FD_ISSET(int fd, *fdset);`
  - Is bit `fd` in `fdset` turned on?