

# COMP/ELEC 429/556

## Introduction to Computer Networks

Reliability

Some slides used with permissions from Edward W.  
Knightly, T. S. Eugene Ng, Ion Stoica, Hui Zhang



# Cannot assume network is reliable

- Goal: Realize the reliable byte stream model
  - Fact: most data networks transmit in granularity of packet – a formatted unit of data
- New goal: Reliably transmit packets and maintain packet order

## Assumptions:

- Packets can get lost or reordered in the network
- Bits in packets can get corrupted on the way

## Solutions?



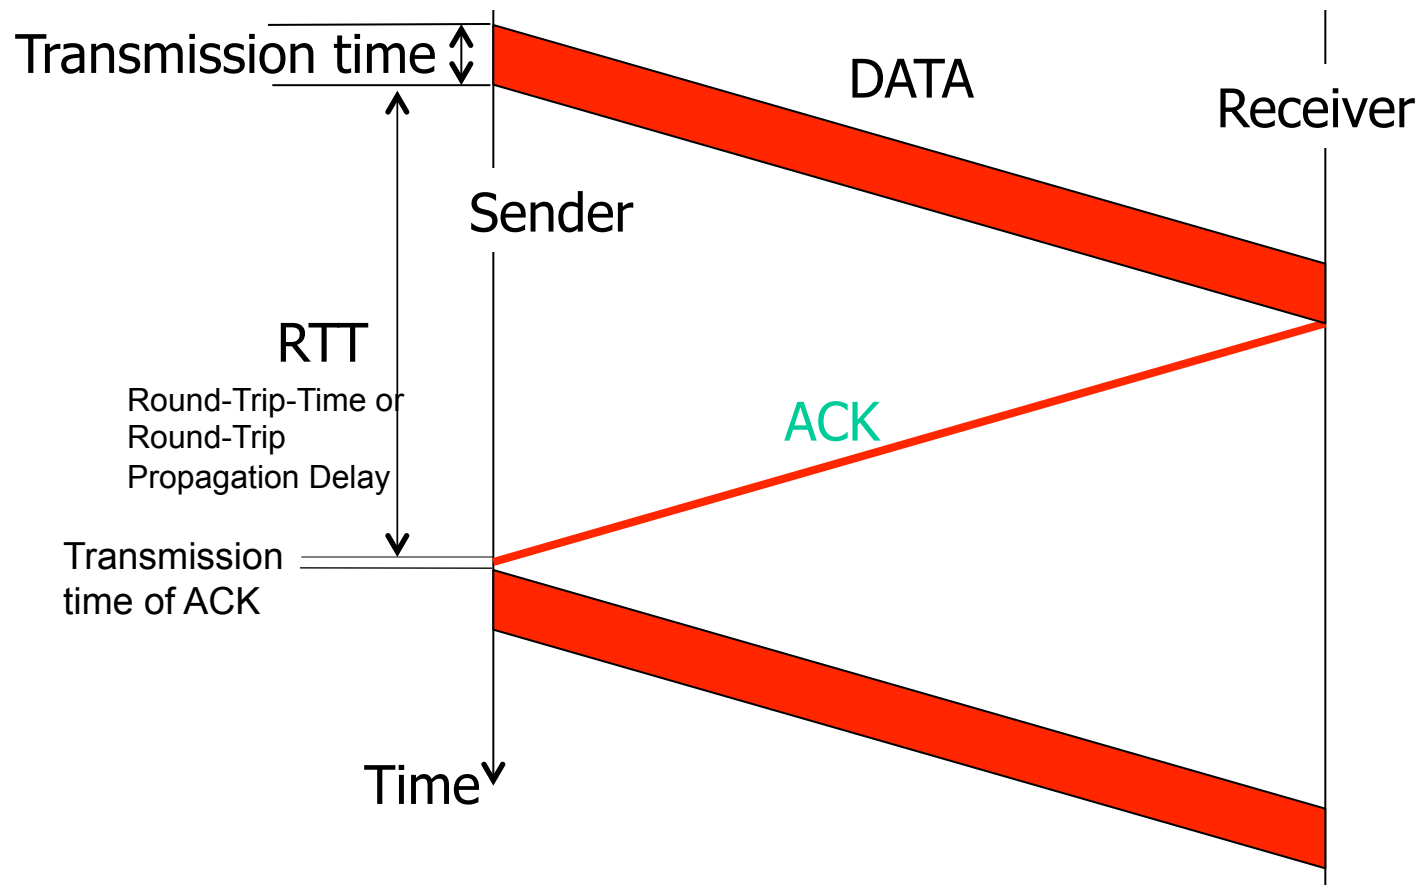
# Retransmission-Based Recovery

- Retransmit lost or corrupted packets
  - Packets are uniquely identified by sequence numbers
  - Correctly received packets are acknowledged
  - Packets not acknowledged are retransmitted
- Key: Do this efficiently
  - Use available bandwidth efficiently
  - Detect loss/corruption and retransmit quickly

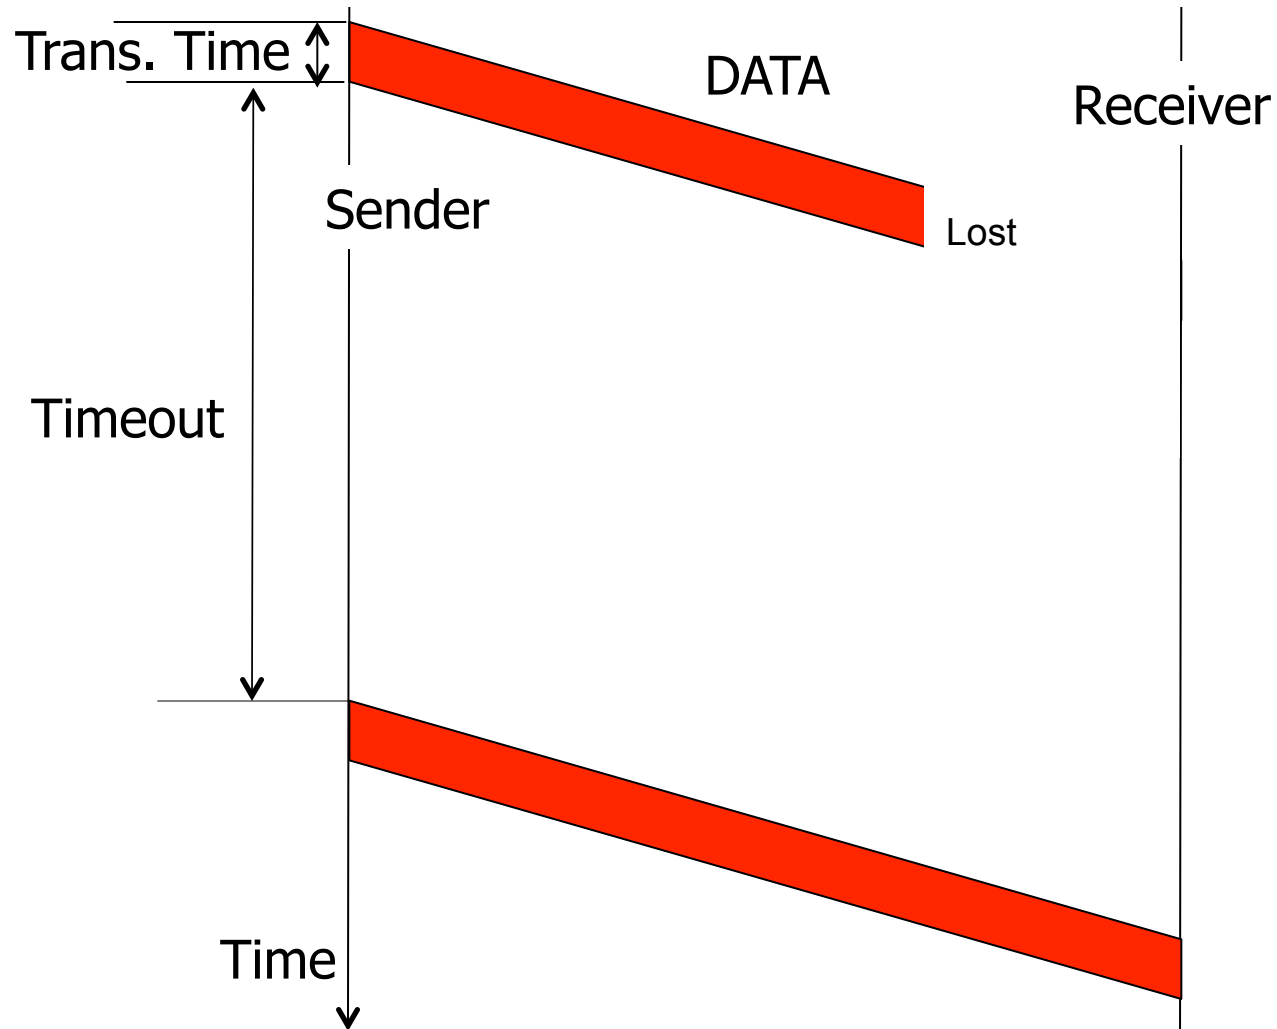


# Naïve Approach: Stop-and-Wait

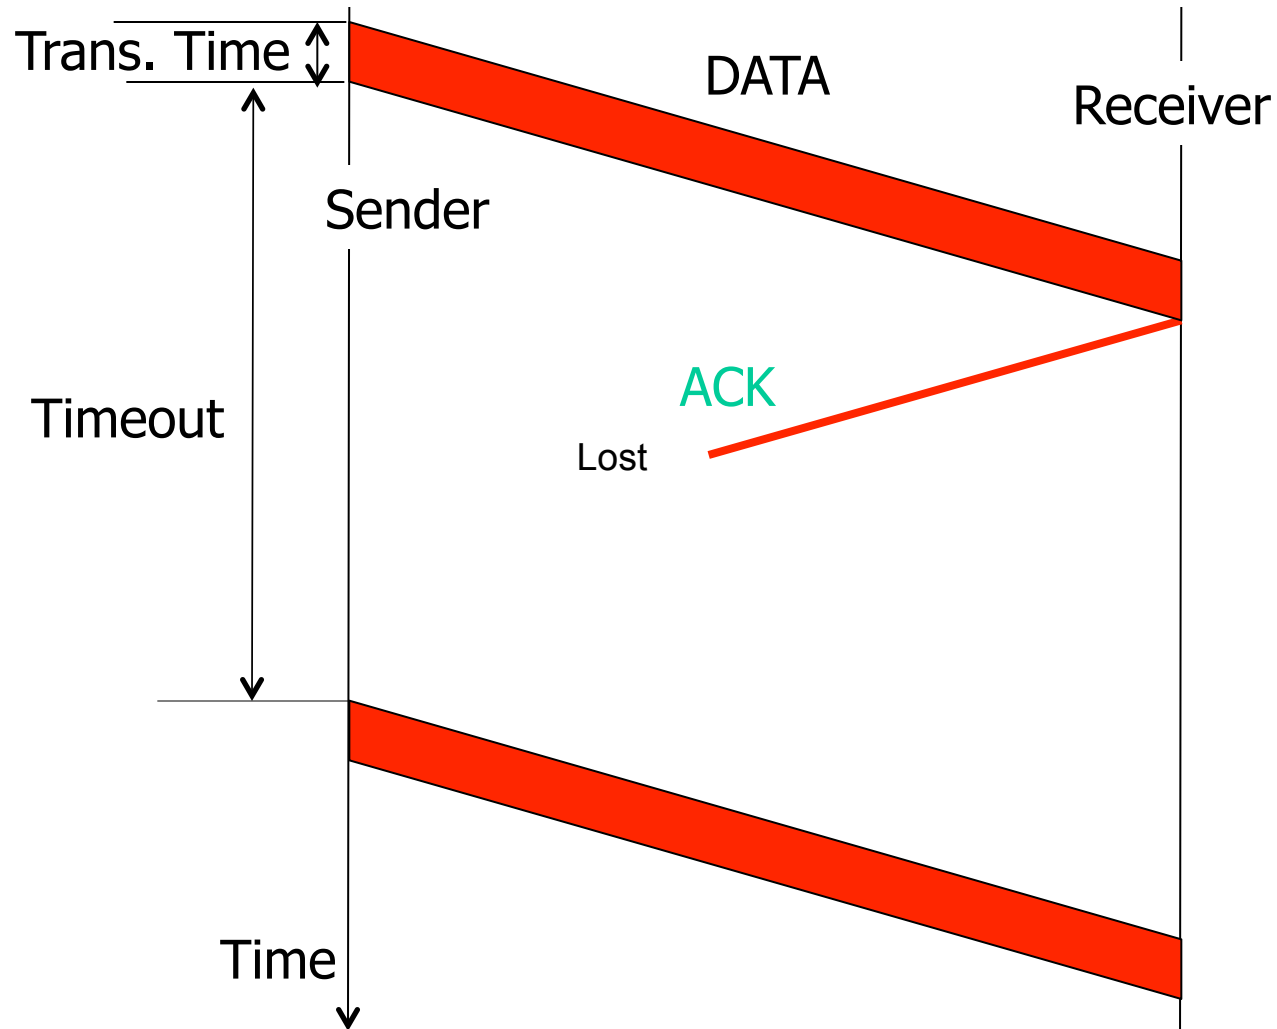
- Send DATA packet; wait for acknowledgement (ACK); repeat
- If timeout, retransmit DATA packet



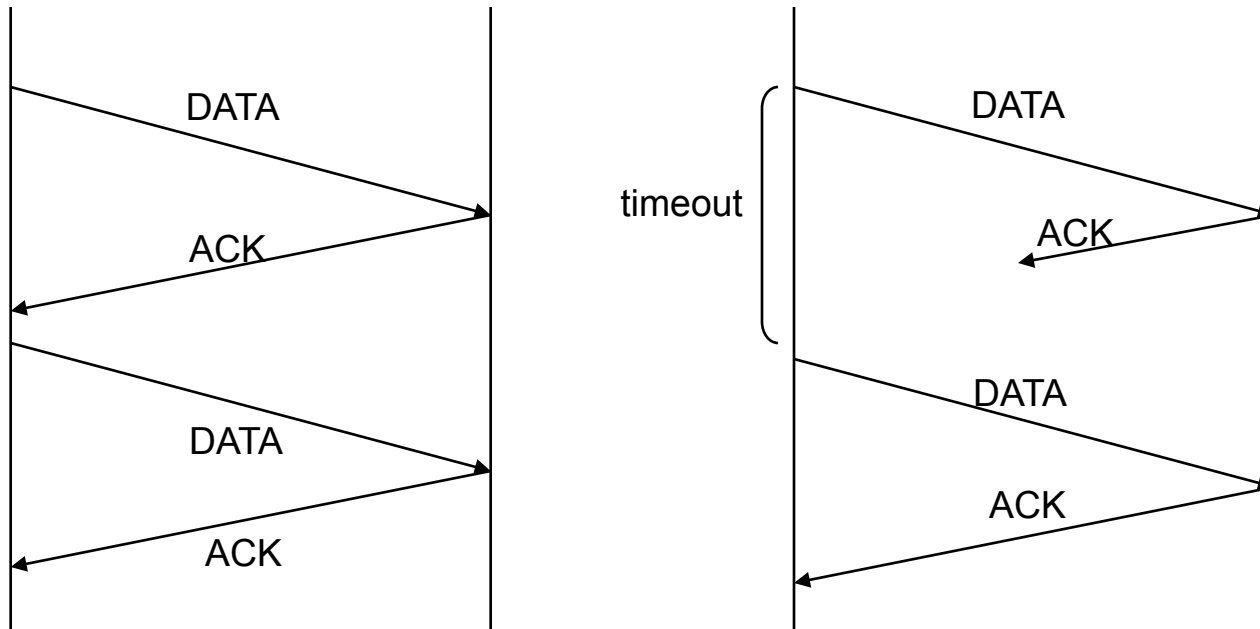
# DATA Loss



# ACK Loss



## How Many Sequence Numbers Needed for Stop-and-Wait?



- 2; need a 1 bit sequence number (i.e. alternate between 0 and 1) to distinguish duplicate packets



## Problem with Stop-and-Wait

- Lots of time wasted in waiting for acknowledgements
- What if you have a 10Gbps link, a data packet size of 1500B (like Ethernet), ACK size of 40B and a propagation delay of 10ms?
- Because you send one packet and wait
  - $\text{Throughput} = 1500 \times 8 \text{bit} / (2 \times 10 \text{ms} + (1500 \text{B} + 40 \text{B}) / 10 \text{Gbps})$   
     $\approx 600 \text{Kbps!}$
  - A utilization of 0.006%


## Better Approach: Sliding Window

- *window* = set of adjacent sequence numbers
- The size of the set is the *window size* (denote it  $n$ )
- Let  $A$  be the last ack'd packet seq # of sender without gap; then window of sender =  $\{A+1, A+2, \dots, A+n\}$



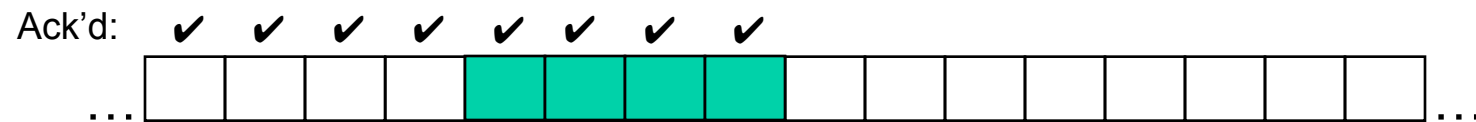
- Sender can send a packet as soon as it is inside its window
- Let  $B$  be the last received packet seq # without gap by receiver; then window of receiver =  $\{B+1, \dots, B+n\}$
- Receiver can accept out of sequence packets if in window

# Sender Sliding Window Example



If receive ACK for packet that falls outside of sliding window, ignore such ACK

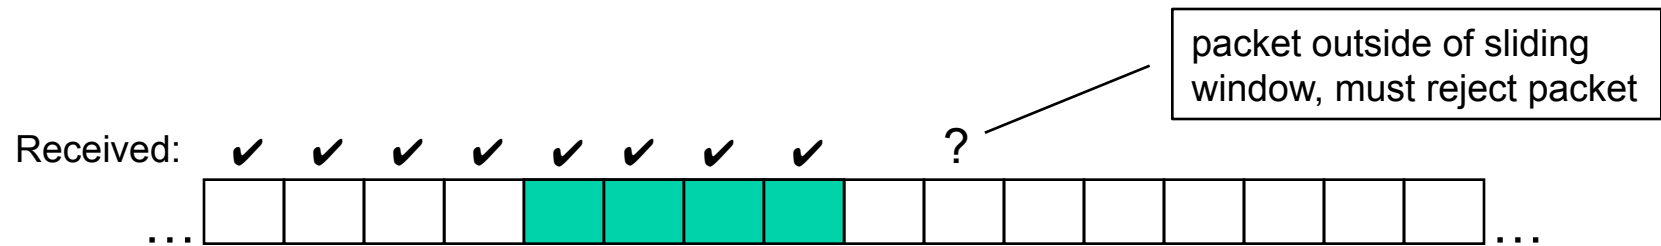
# Sender Sliding Window Example



# Sender Sliding Window Example



# Receiver Sliding Window Example



## Basic Timeout and Acknowledgement


- Every packet  $k$  transmitted is associated with a timeout, denoted  $\text{timeout}(k)$
- Basic acknowledgement scheme
  - Upon receiving a packet from sender, send an ack for  $k$ , where  $k$  is the newest packet seq # such that all packets with seq numbers older than or equal to  $k$  have been accepted



- Suppose packet B is received but A is missing, ack for A-1 is sent. If C is received next, ack for A-1 is sent again. If D is received next, ack for A-1 is sent again. If A is received next, an ack for D is sent, and the receiver window slides
- If by  $\text{timeout}(k)$ , ack for  $k$  or larger has not yet been received, the sender retransmits packet  $k$


## Example with Loss

Window size = 3 packets

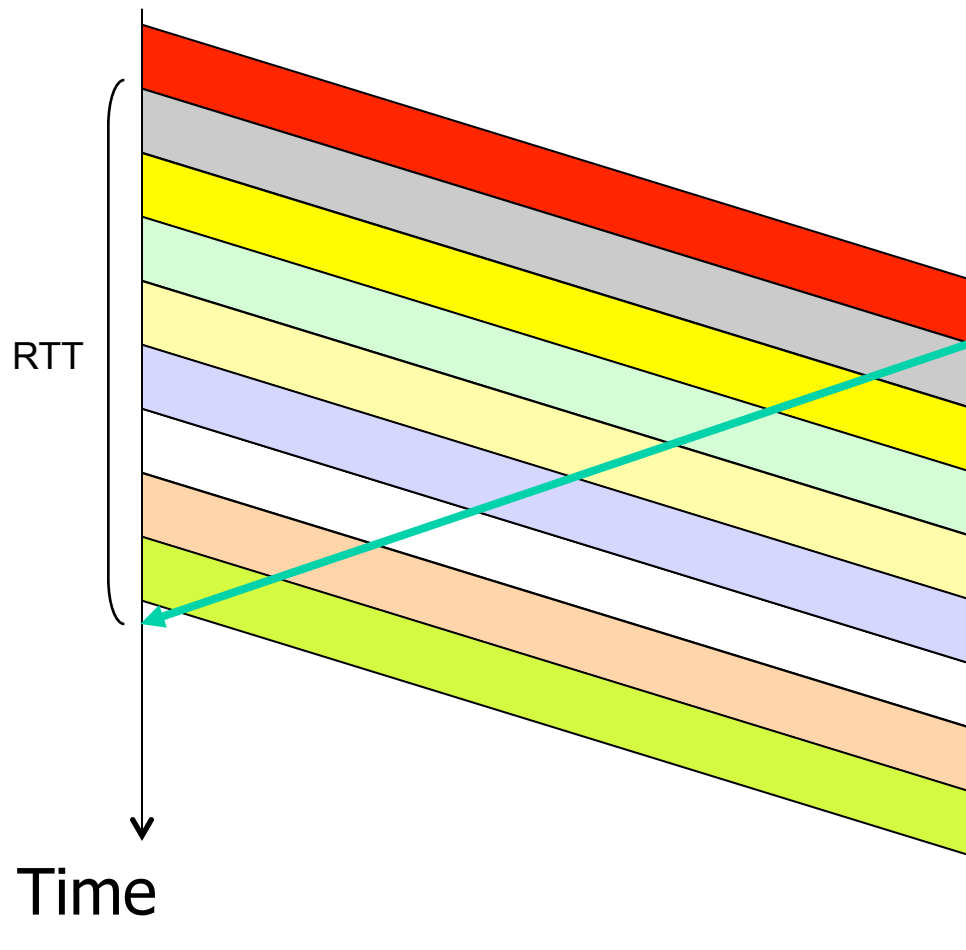




# Packets within Sender Window Can Be Sent



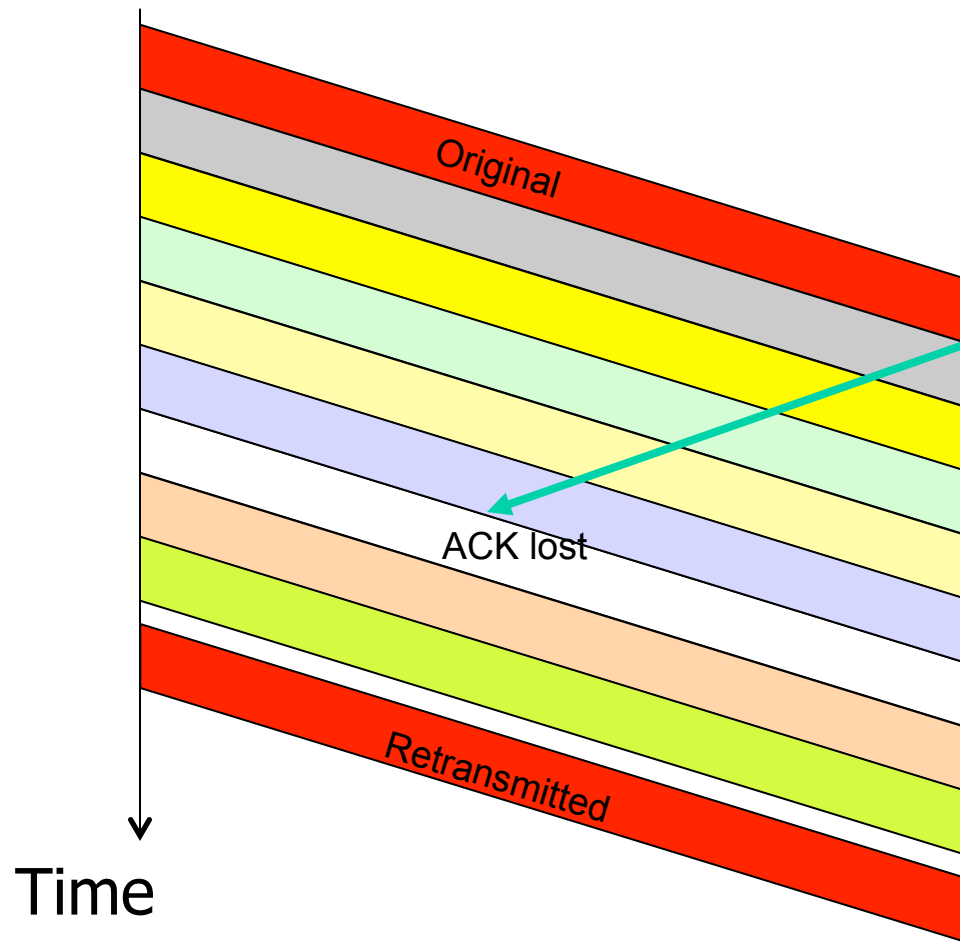
# Efficiency



$n = 9$ , i.e. 9 packets in one RTT

Can fully use available bandwidth if  $n$  is large enough

# How Many Unique Sequence Numbers Needed?



n unique sequence numbers not enough to differentiate the retransmission of a packet in current window and a new transmission from the next window

Need  $2*n$  unique sequence numbers

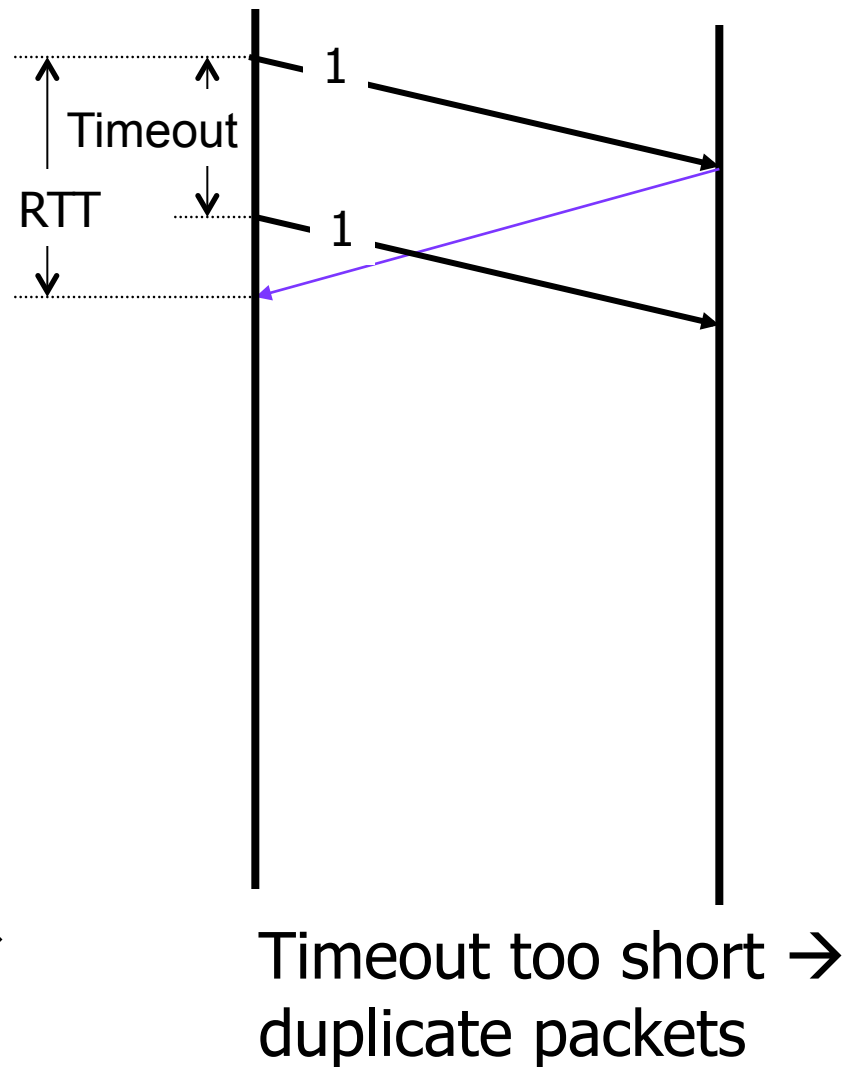

## Observations

- With sliding windows, it is possible to fully utilize a link, provided the window size is large enough  
Throughput proportional to  $(n/RTT)$ 
  - Stop & Wait is like  $n = 1$
- But how to figure out the right window size to use???
  - Should be no larger than what's sustainable by the network and the sender/receiver machines
- Sender has to buffer all unacknowledged packets, because they may require retransmission
- Receiver may accept out-of-order packets within window, and delay delivery to application until the missing in-order packets are received

## Setting Timeout Value

- The sender needs to pick a retransmission timeout value in order to decide when to retransmit a packet that may have been lost
- How to pick the timeout value?
  - **Too short**: may retransmit before data or ACK has arrived, creating duplicates
  - **Too long**: if a packet is lost, will take a long time to recover (inefficient)

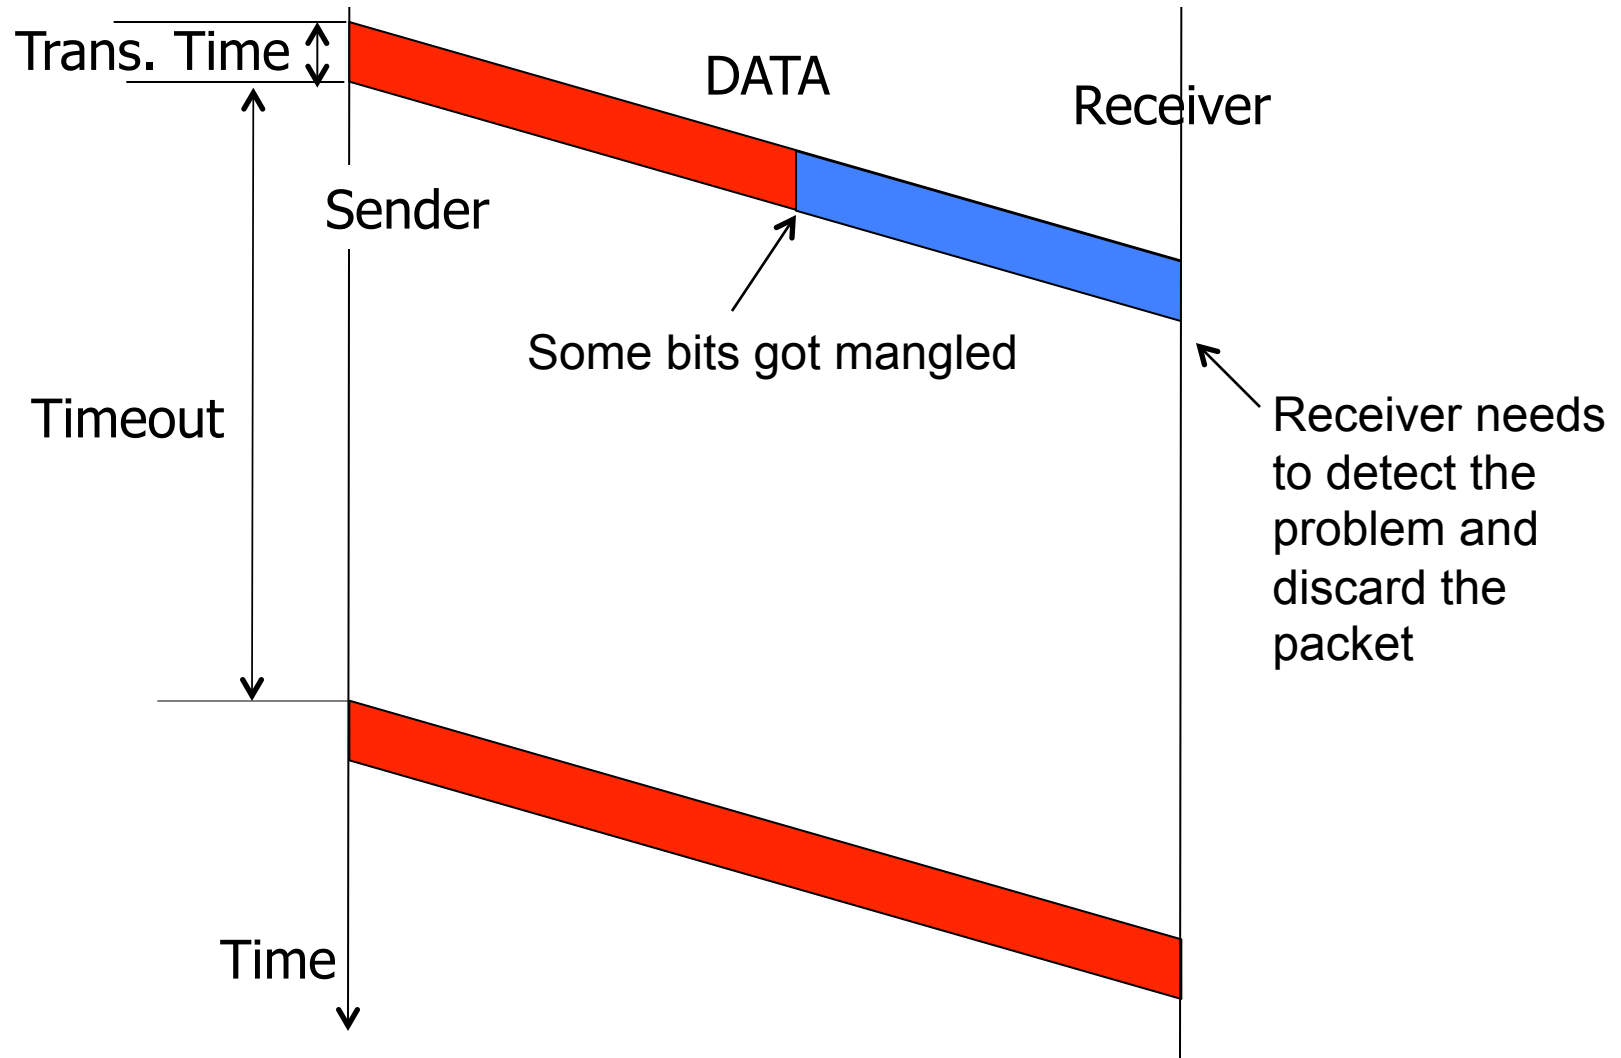
# Timing Illustration



## Adaptive Timeout Value

- The amount of time the sender should wait is related to the round-trip time (RTT) between the sender and receiver
- RTT not known a priori and may change, so a protocol should measure RTT and adapt timeout value accordingly
- The timeout value should be larger than the RTT

# Data Bit Error/Corruption





# Error Detection: Naïve Approach

- Send a message twice
- Compare two copies at the receiver
  - If different, some errors exist
- How many bits of error can you detect?
- What is the overhead?



# Error Detection

- Problem: detect bit errors in packets
- Solution: add **extra** bits to each packet
- Goals:
  - Small overhead, i.e., small number of redundancy bits
  - Large number of bit error patterns that can be detected



## What is the Relationship between the red bit and the black bits?

e.g. 1001111 1

e.g. 1111011 0

e.g. 1011111 0

e.g. 1010001 1

e.g. 0100101 1

e.g. 1110111 0

# Parity – Illustrates Idea Behind Error Detection

- Even parity
  - Add a parity bit to 7 bits of data to make an even number of 1's

0110100	1
1011010	0

- How many bits of error can be detected by a parity bit?
- What's the overhead?

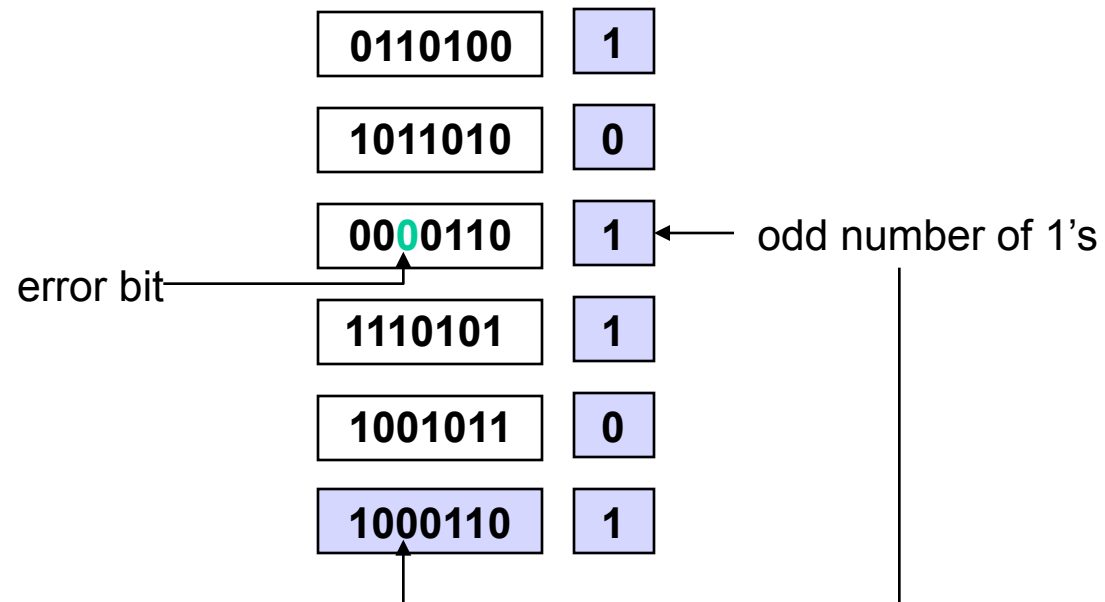
# Two-dimensional Parity

- Add one extra bit to a 7-bit code such that the number of 1's in the resulting 8 bits is even (for even parity, and odd for odd parity)
- Add a parity byte for the packet
- Example: five 7-bit character packet, even parity

0110100	1
1011010	0
0010110	1
1110101	1
1001011	0
1000110	1

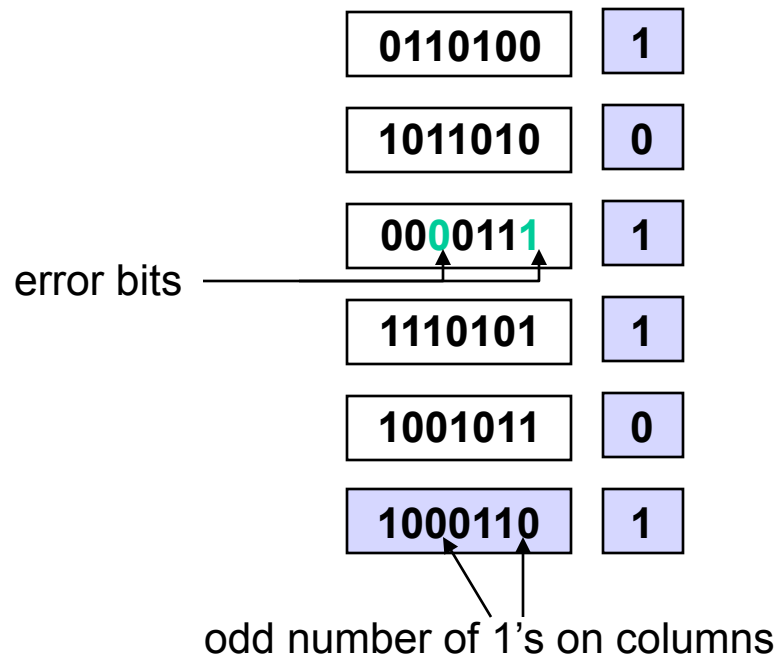
# How Many Errors Can you Detect?

- All 1-bit errors
- Example:



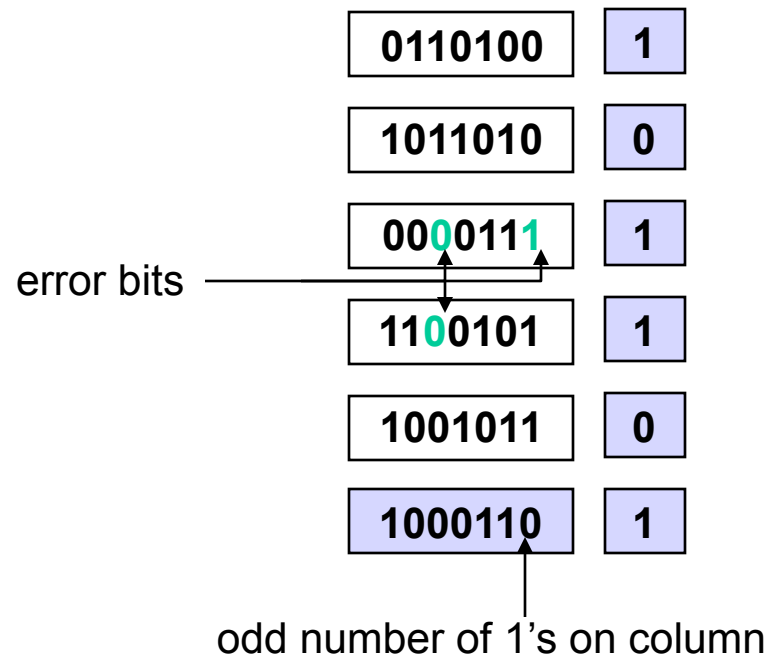
# How Many Errors Can you Detect?

- All 2-bit errors
- Example:



# How Many Errors Can you Detect?

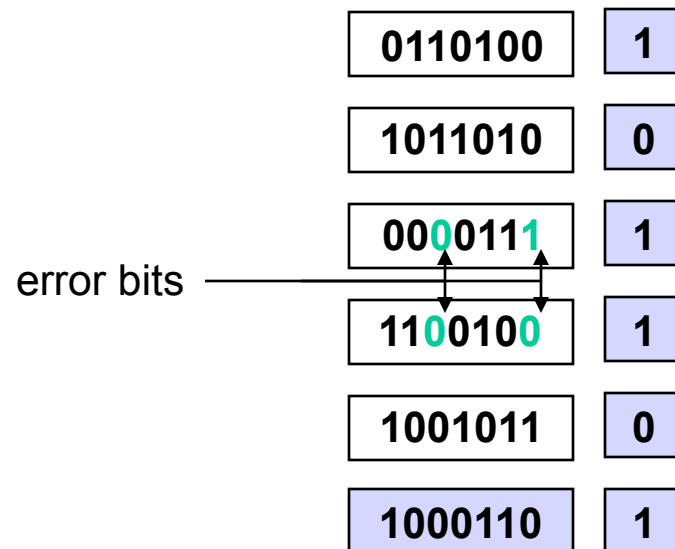
- All 3-bit errors
- Example:





# How Many Errors Can you Detect?


- Example of 4-bit error that is **not** detected:



# Error Detection in Network Protocols

- Parity is used in PCI buses, cache/memory, etc.
- Parity is seldom used in network protocols
- Ethernet uses cyclic redundancy check
- Internet protocols use 1's complement sum

# Cyclic Redundancy Check (CRC)



## Arithmetic Modulo 2

- Like binary arithmetic but without borrowing/carrying from/to adjacent bits
- Examples:

$$\begin{array}{r} 101 + \\ 010 \\ \hline 111 \end{array} \quad \begin{array}{r} 101 + \\ 001 \\ \hline 100 \end{array} \quad \begin{array}{r} 1011 + \\ 0111 \\ \hline 1100 \end{array}$$

$$\begin{array}{r} 101 - \\ 010 \\ \hline 111 \end{array} \quad \begin{array}{r} 101 - \\ 001 \\ \hline 100 \end{array} \quad \begin{array}{r} 1011 - \\ 0111 \\ \hline 1100 \end{array}$$

- Addition and subtraction in binary arithmetic modulo 2 is equivalent to XOR

a	b	$a \otimes b$
0	0	0
0	1	1
1	0	1
1	1	0

# Some Polynomial Arithmetic Modulo 2

## Properties

- Subtracting/adding  $C(x)$  from/to  $B(x)$  modulo 2 is equivalent to performing an XOR on each pair of matching coefficients of  $C(x)$  and  $B(x)$

– E.g.:

$$\begin{array}{rcl} B(x) & = & x^7 + x^5 + x^3 + x^2 + x^0 \quad (10101101) \\ C(x) & = & x^3 + x^1 + x^0 \quad (00001011) \\ \hline B(x) - C(x) & = & x^7 + x^5 + x^2 + x^1 \quad (10100110) \end{array}$$

# Cyclic Redundancy Check (CRC)

- Represent a n-bit message as an (n-1) degree polynomial  $M(x)$ 
  - E.g., 10101101  $\rightarrow M(x) = x^7 + x^5 + x^3 + x^2 + x^0$
- Choose a divisor k-degree polynomial  $C(x)$
- Compute remainder  $R(x)$  of  $M(x)*x^k / C(x)$ , i.e.,  
 $M(x)*x^k = A(x)*C(x) + R(x)$ , where  $\text{degree}(R(x)) < k$
- Let  
 $T(x) = M(x)*x^k - R(x) = A(x)*C(x)$
- Then
  - $T(x)$  is divisible by  $C(x)$
  - First n coefficients of  $T(x)$  represent  $M(x)$

# Cyclic Redundancy Check (CRC)

- Sender:
  - Compute and send  $T(x)$ , i.e., the coefficients of  $T(x)$
- Receiver:
  - Let  $T'(x)$  be the  $(n-1+k)$ -degree polynomial generated from the received message
  - If  $C(x)$  divides  $T'(x) \rightarrow$  assume no errors; otherwise errors
- Note: all computations are modulo 2

## Example (Sender Operation)

- Send packet 110111; choose  $C(x) = 101$ 
  - $k = 2$ ,  $M(x) \cdot x^k \rightarrow 11011100$
- Compute the remainder  $R(x)$  of  $M(x) \cdot x^k / C(x)$

$$\begin{array}{r}
 101 \overline{) 11011100} \\
 \underline{101} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 111 \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 \underline{101} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 101 \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 \underline{101} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 100 \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 \underline{101} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 1 \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \leftarrow R(x)
 \end{array}$$

- Compute  $T(x) = M(x) \cdot x^k - R(x) \rightarrow 11011100 \text{ xor } 1 = 11011101$
- Send  $T(x)$



## Example (Receiver Operation)

- Assume  $T'(x) = 11011101$ 
  - $C(x)$  divides  $T'(x) \rightarrow$  no errors
- Assume  $T'(x) = 11001101$ 
  - Remainder  $R'(x) = 1 \rightarrow$  error!

$$\begin{array}{r}
 101 \overline{) 11001101} \\
 \underline{101} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 110 \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 \underline{101} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 111 \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 \underline{101} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 101 \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 \underline{101} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 101 \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 \underline{101} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \\
 1 \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00} \phantom{00}
 \end{array}$$

1  $\leftarrow R'(x)$

- Note: an error is **not** detected iff  $C(x)$  divides  $T'(x) - T(x)$

## CRC Properties

- Detect all single-bit errors if coefficients of  $x^k$  and  $x^0$  of  $C(x)$  are one
- Detect all double-bit errors, if  $C(x)$  has a factor with at least three terms
- Detect all number of odd errors, if  $C(x)$  contains factor  $(x+1)$
- Detect all burst of errors smaller than  $k$  bits
- See Peterson & Davie Table 2.3 for some commonly used CRC polynomials

## 1's Complement Checksum



- Sender: add all words of data and append the result (checksum) to the data
- Receiver: add all words of received data and compare the result with the checksum

# 1's Complement Number Representation

- Negative number  $-x$  is  $x$  with all bits inverted
- When two numbers are added, the carry-on is added to the result
- Example:  $-15 + 16$ ; assume 8-bit representation

15 = 00001111  $\rightarrow$  -15 = 11110000

+

16 = 00010000

---

1 00000000

+ (dashed line)

1

---

00000001

-15 + 16 = 1

The diagram illustrates the 1's complement addition of -15 and 16. It shows the conversion of 15 to its 1's complement representation -15 (11110000). Then, 16 (00010000) is added to -15. The result is 00000001, with a carry-in of 1 from the 9th bit position. The final result is 1, which is the sum of -15 and 16.

# Software Implementation

```
u_short cksum(u_short *buf, int count)
{
    register u_long sum = 0;
    while (count--)
    {
        sum += *buf++;
        if (sum & 0xFFFF0000)
        {
            /* carry occurred, so wrap around */
            sum &= 0xFFFF;
            sum++;
        }
    }
    return ~(sum & 0xFFFF);
}
```

# Properties

- How many bits of error can 1's complement sum detect?
- What's the overhead?
- Why use this algorithm?
  - Need to understand how different technologies coordinate with each other
  - We'll start this journey in the next part of the course

