

COMP/ELEC 429/556

Introduction to Computer Networks


Principles of Congestion Control

Some slides used with permissions from Edward W.
Knightly, T. S. Eugene Ng, Ion Stoica, Hui Zhang




What is Congestion?

- The load placed on the network is higher than the capacity of the network
 - Not surprising: independent senders place load on network
- Results in packet loss: routers have no choice
 - Can only buffer finite amount of data




How Fast to Send? What's at Stake?

- Send too slow: link sits idle
 - wastes time
- Send too fast: link is kept busy but....
 - queue builds up in router buffer (delay)
 - overflow buffers in routers (loss)
 - Many retransmissions, many losses
 - Network goodput (throughput of useful data) goes down
 - “Congestion collapse”




Abstract View




- We ignore internal structure of network and model it as having a single bottleneck link

Problem 1: Single Flow, Fixed Bandwidth



- Adjust rate to match bottleneck bandwidth
 - without any *a priori* knowledge
 - could be 40 Gbps link, could be a 32 Kbps link

Problem 2: Single Flow, Varying Bandwidth




- Adjust rate to match instantaneous bandwidth
- Bottleneck can change because of a routing change

Problem 3: Multiple Flows

Two Issues:


- Adjust total sending rate to match bottleneck bandwidth
- Allocation of bandwidth between flows



General Approaches

- Reservation
 - pre-arrange bandwidth allocations
 - requires negotiation before sending packets
 - requires router support

Sliding Window



Window size $n = 9$, i.e.
9 packets in one RTT

In general, sending rate
proportional to n/RTT

General Approaches (cont'd)

- Dynamic sending rate adjustment
 - Every sender probe network to test level of congestion
 - speed up when no congestion
 - slow down when congestion
 - suboptimal, messy dynamics, but simple to implement
 - requires no router support
- Distributed coordination problem!



Sliding Window Congestion Control

- Sender has a send window
 - controls amount of unacknowledged data in transit
- Sending rate proportional to: $\text{Send window size} / \text{RTT}$
- Vary send window size to control sending rate

Two Basic Components

- Detecting congestion
- Rate adjustment algorithm (change window size)
 - depends on congestion or not

Detecting Congestion

- Packet dropping is a plausible sign of congestion
 - delay-based methods are hard and risky
- How do you detect packet drops? ACKs
 - ACKs signal receipt of data
 - ACK denotes last contiguous byte received
- Two signs of packet drops
 - No ACK after certain time interval: time-out
 - Several duplicate ACKs for the same sequence number
- This heuristic may not work well for wireless networks, why?
 - Think whether packet drops are always due to congestion


Rate Adjustment

- Basic idea:
 - Upon receipt of ACK (of new data): increase rate
 - Data successfully delivered, perhaps can send faster
 - Upon detection of loss: decrease rate
- But how much increase/decrease should be applied?
 - What outcomes do we want?
- For simplicity, restrict to “additive” and “multiplicative” increase/decrease
 - “additive” results in linearly change
 - “multiplicative” results in exponential change



Fairness & Efficiency

Two competing sessions:


- Additive increase (AI) gives slope of 1, as throughput increases
- multiplicative decrease (MD) decreases throughput proportionally




AIMD





AIMD Sharing Dynamics



- No congestion \rightarrow rate increases by one packet/RTT every RTT
- Congestion \rightarrow decrease rate by factor 2




AIAD




Limit rates:
x and y depend
on initial
values

AIAD Sharing Dynamics




- No congestion \rightarrow x increases by one packet/RTT every RTT
- Congestion \rightarrow decrease x by 1



AIMD Model


- Analyze the steady state throughput as a function of
 - RTT
 - Loss probability
- Assumptions
 - Each packet dropped with *iid* probability p
- Methodology: analyze “average” cycle in steady state
 - How many packets are transmitted per cycle?
 - What is the duration of a cycle?

Cycles in Steady State



- Denote W as the maximum achieved window
- What is the slope of the line?
- What are the key values on the time axis?

Cycle Analysis



W increase by 1 per RTT

$$\text{pkts xmitted/cycle} = \text{area} = \left(\frac{W}{2}\right)^2 + \frac{1}{2}\left(\frac{W}{2}\right)^2 = \frac{3}{8}W^2$$

Throughput

$$\text{throughput} = \frac{\text{pkts xmitted/cycle}}{\text{time/cycle}} = \frac{\frac{3}{8}W^2}{RTT\left(\frac{W}{2}\right)}$$

- What is W as a function of p ?
How long does a cycle last until a drop?

Cycle Length

Let α be the index of the lost packet that ends a cycle

$$\begin{aligned} P(\alpha = k) &= P(k-1 \text{ pkts not lost, } k\text{th pkt lost}) \\ &= (1-p)^{k-1} p \end{aligned}$$

$$\Rightarrow E(\alpha) = \sum_{k=1}^{\infty} k(1-p)^{k-1} p = \frac{1}{p}$$

$$\Rightarrow \frac{1}{p} = \frac{3}{8} W^2 \quad \Rightarrow W = \sqrt{\frac{8}{3p}}$$

AIMD Model

$$\text{throughput } T(p) = \frac{\frac{1}{p}}{RTT \cdot \frac{1}{2} \sqrt{\frac{8}{3p}}} = \frac{1}{RTT \sqrt{\frac{2}{3}} p}$$

- Note role of RTT. Is it “fair”?
- A “macroscopic” model