

COMP/ELEC 429/556

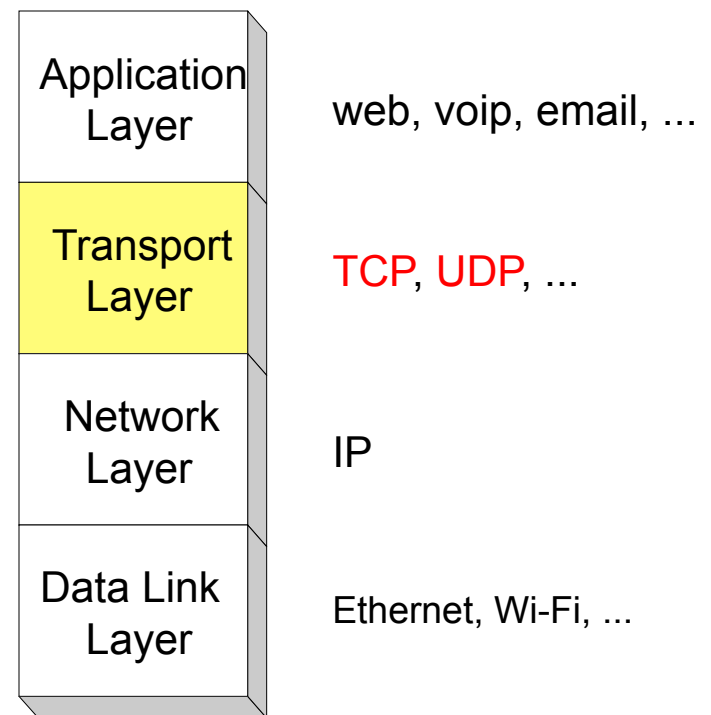
Introduction to Computer Networks

The TCP Protocol

Some slides used with permissions from Edward W.
Knightly, T. S. Eugene Ng, Ion Stoica, Hui Zhang


Transport Layer

- Purpose 1: Demultiplexing of data streams to different application processes
- Purpose 2: Provide value-added services that many applications want
 - Recall network layer in Internet provides a “Best-effort” service only, transport layer can add value to that
 - Application may want reliability, etc
 - No need to reinvent the wheel each time you write a new application




A very simple transport protocol: User Datagram Protocol (UDP)

- Connectionless datagram
 - Socket: SOCK_DGRAM
- Port number used for demultiplexing
 - port numbers = connection/application endpoint
- Adds end-to-end error checking through optional checksum
 - some protection against data corruption errors between source and destination (links, switches/routers, bus)
 - does not protect against packet loss, duplication or reordering



Using Transport Layer Port Number to Demultiplex traffic

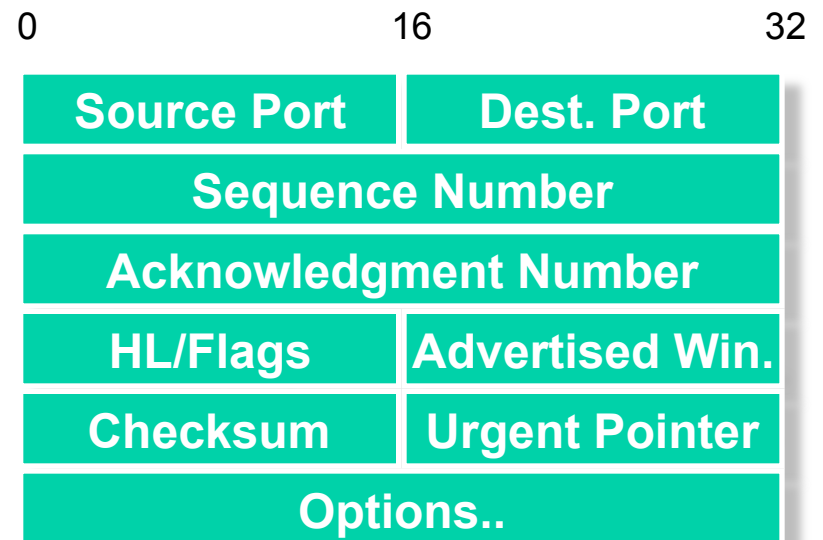


Usages of UDP

- Custom protocols/applications can be implemented on top of UDP
 - use the port addressing provided by UDP
 - implement specialized reliability, flow control, ordering, congestion control as the app sees fit
- Examples:
 - remote procedure call
 - multimedia streaming (real time protocol)
 - cluster computing communication libraries

Transmission Control Protocol (TCP)

- Reliable bidirectional in-order byte stream
 - Socket: SOCK_STREAM
- Connections established & torn down
- Multiplexing/ demultiplexing
 - Ports at both ends
- Error control
 - Users see correct, ordered byte sequences
- End-to-end flow control
 - Avoid overwhelming receiver at each end
- Congestion control
 - Avoid creating traffic jams within network



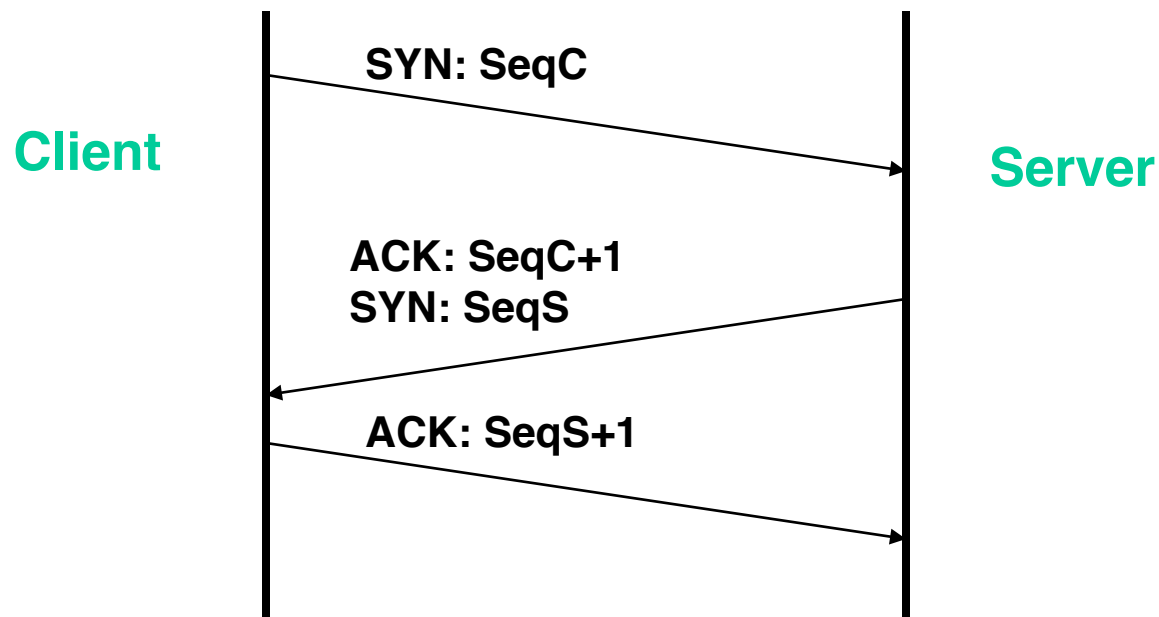
Connection Setup

- Why need connection setup?
- Mainly to agree on starting sequence numbers
 - Starting sequence number is randomly chosen
 - Reason: to reduce the chance that sequence numbers of old and new connections from overlapping

Important TCP Flags

- SYN: Synchronize
 - Used when setting up connection
- FIN: Finish
 - Used when tearing down connection
- ACK
 - Acknowledging received data

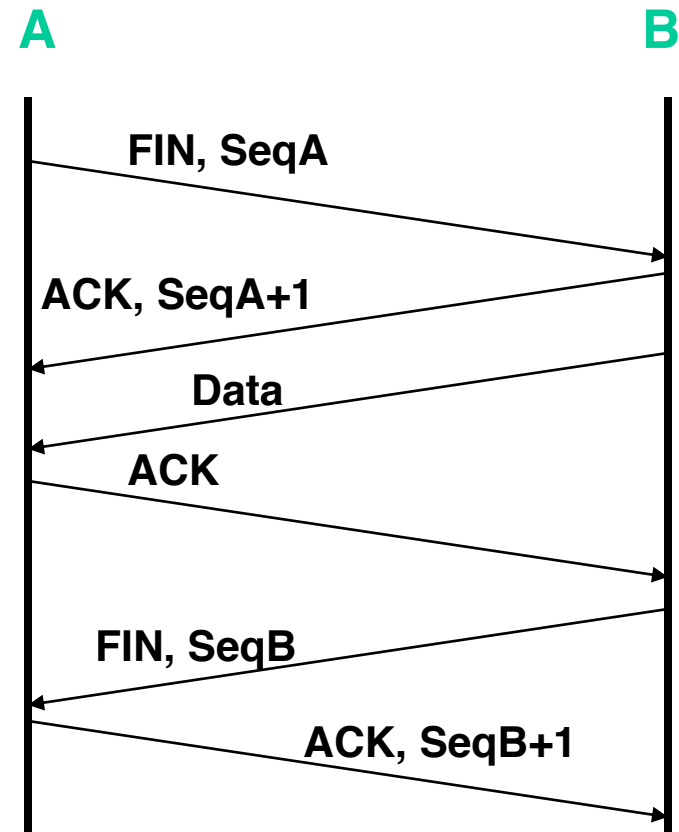
Establishing Connection



- Three-Way Handshake
 - Each side notifies other of starting sequence number it will use for sending
 - Each side acknowledges other's sequence number
 - SYN-ACK: Acknowledge sequence number + 1
 - Can combine second SYN with first ACK


Tearing Down Connection

- Either Side Can Initiate Tear Down
 - Send FIN signal
 - “I’m not going to send any more data”
- Other Side Can Continue Sending Data
 - Half open connection
 - Must continue to acknowledge
- Acknowledging FIN
 - Acknowledge last sequence number + 1




Sequence Number Space

- Each byte in byte stream is numbered
 - 32 bit value
 - Wraps around
 - Initial values selected at start up time
- TCP breaks up the byte stream into segments
 - Each segment transmitted by a packet
 - Limited by the Maximum Segment Size
 - Set to prevent packet fragmentation
- Each segment has a sequence number
 - Indicates where it fits in the byte stream



Bidirectional Communication



- Each Side of Connection can Send *and* Receive
- What this Means
 - Maintain different sequence numbers for each direction
 - A single packet can contain new data for one direction, plus acknowledgement for other, may also contain only data or only acknowledgement
- When there is a loss, e.g. seq # 2000, 2002, 2003, 2004 are received, TCP receiver acks 2001, 2001, 2001, 2001
 - Called duplicate ACKs
 - There are TCP variants that don't do this – beyond our scope

Sequence Numbers

- 32 Bits, Unsigned
- Why So Big?
 - For sliding window, must have
 $|\text{Sequence Space}| > 2 * |\text{Sending Window}|$
 - Sending window size of basic TCP is at most 2^{16} bytes
 - $2^{32} > 2 * 2^{16}$; no problem
 - Also, want to guard against stray packets
 - With IP, assume packets have maximum segment lifetime (MSL) of 120s
 - i.e. can linger in network for upto 120s
 - Sequence number would wrap around in this time at 286Mbps

Error Control

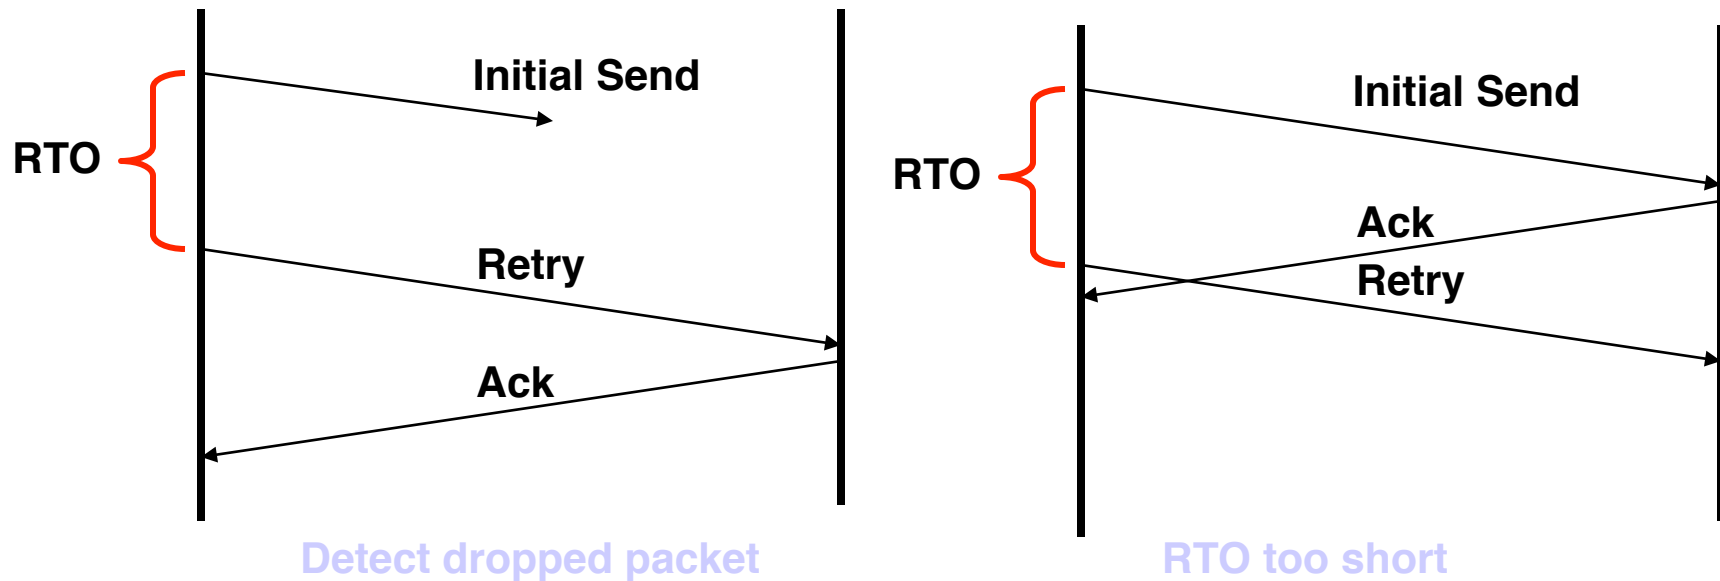
- Checksum provides some end-to-end error protection
- Sequence numbers detect packet sequencing problems:
 - Duplicate: ignore
 - Reordered: reorder or drop
 - Lost: retransmit
- Lost segments retransmitted by sender
 - Use time out to detect lack of acknowledgment
 - Need estimate of the roundtrip time to set timeout
- Retransmission requires that sender keep copy of the data
 - Copy is discarded when ack is received

TCP Must Operate Over Any Internet Path

- Retransmission time-out should be set based on round-trip delay
- But round-trip delay different for each path!
- Must estimate RTT dynamically




Setting Retransmission Timeout (RTO)



- Time between sending & resending segment
- Challenge
 - Too long: Add latency to communication when packets dropped
 - Too short: Send too many duplicate packets
 - General principle: Must be > 1 Round Trip Time (RTT)

Round-trip Time Estimation

- Every Data/Ack pair gives new RTT estimate

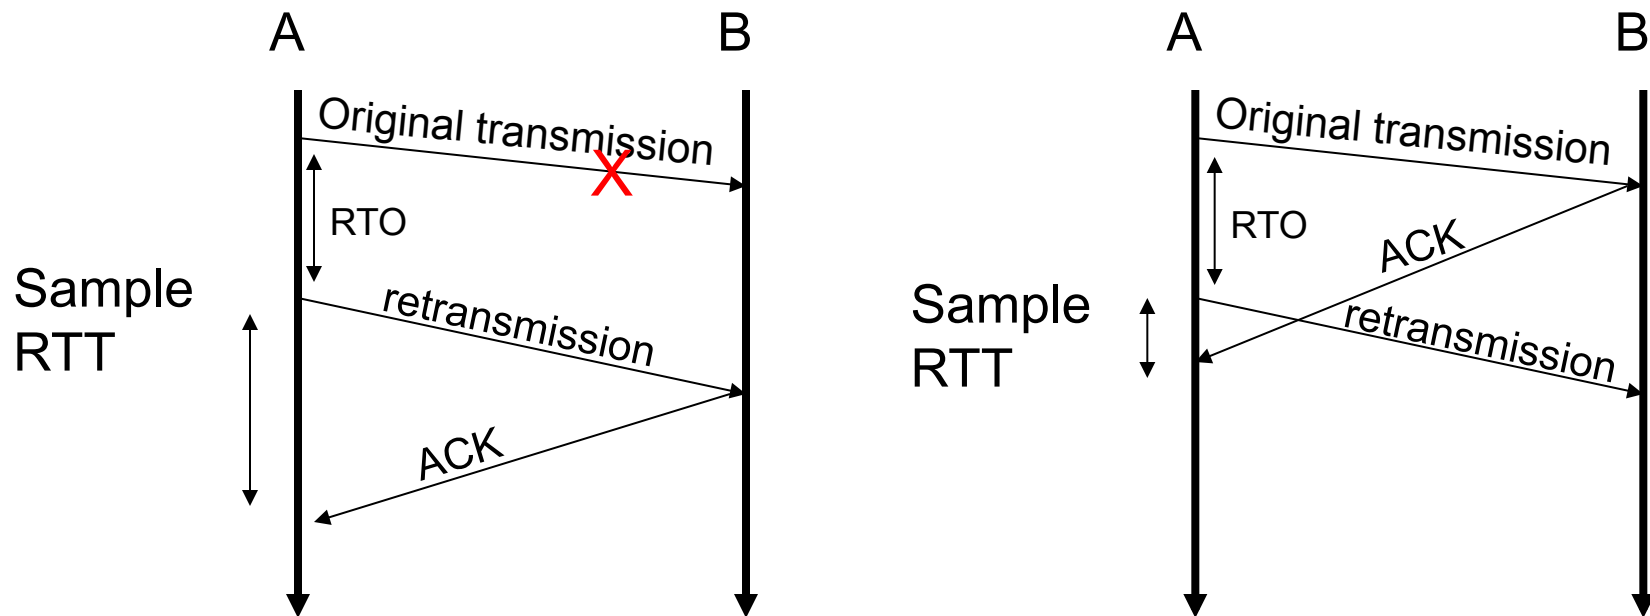


- Can Get Lots of Short-Term Fluctuations
 - How to address this problem?

Exponential Smoothing Technique

- Round trip times estimated as a moving average:
 - Smoothed RTT = α (Smoothed RTT) + $(1 - \alpha)$ (new RTT sample)
 - Recommended value for α : 0.8 - 0.9
- Retransmission timeout (RTO) is a function of
 - Smoothed RTT (SRTT)
 - RTT variation (RTTVAR)
- $RTO = SRTT + 4 * RTTVAR$
 - Details in RFC 6298

RTT Sample Ambiguity

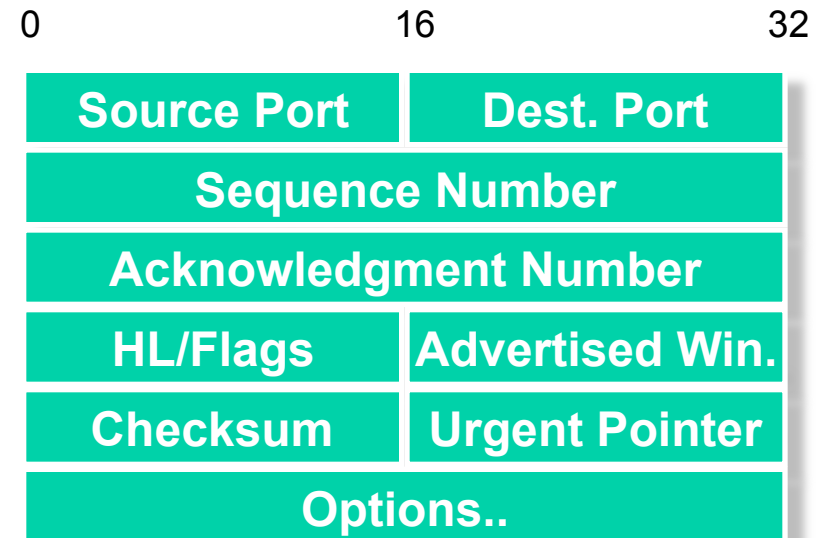


- Ignore sample for segment that has been retransmitted

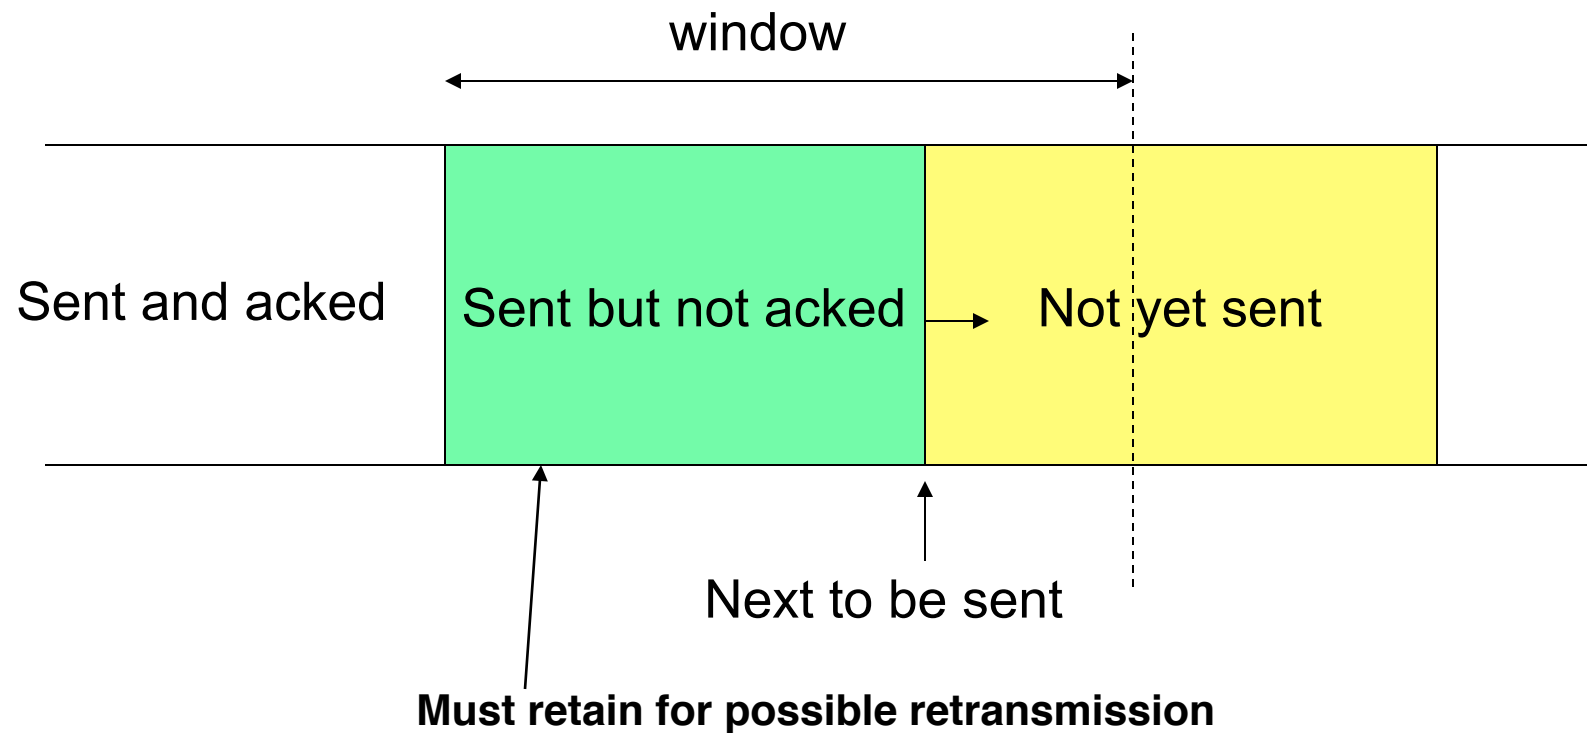
TCP Speed Control

- Sliding window protocol
 - For window size n , can send up to n bytes without receiving new acknowledgement
 - Speed proportional to n/RTT
 - When the data are acknowledged then the window slides forward

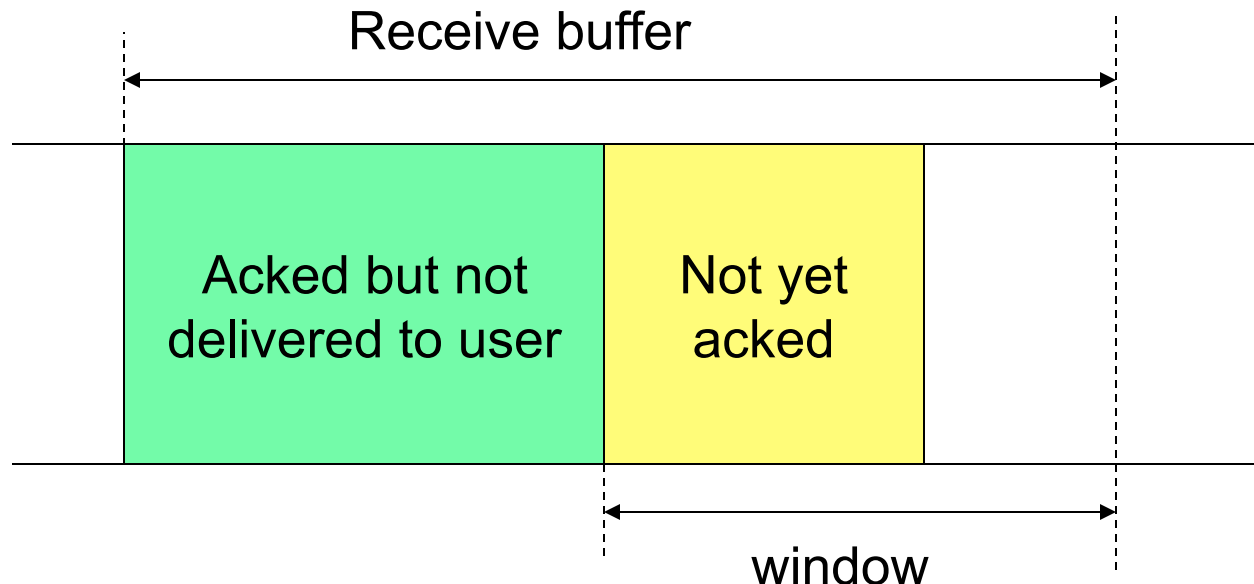
- Send window size set to **minimum (advertised window, congestion window)**



Window Flow Control: Send Side



Window Flow Control: Receive Side



- TCP receiver can delete acknowledged data only after the data has been delivered to the application
- So, depending on how fast the application is reading the data, the receiver's window size may change!!!

Solution

- Receiver tells sender the current advertised window size in every packet it transmits to the sender
- Sender uses this current advertised window size as an upper bound
 - send window size = minimum (advertised window, congestion window)
- Advertised window size is continuously changing
- Can go to zero!
 - Sender not allowed to send anything!

Setting Congestion Window

Send window size = minimum (advertised window, **congestion** window)

Phases of TCP congestion control

1. Slow start (getting to equilibrium)
 - Want to find this very very fast and not waste time
2. Congestion Avoidance
 - Additive increase - gradually probing for additional bandwidth
 - Multiplicative decrease – decreasing congestion window upon loss/timeout

Variables Used in Implementation

- **Congestion Window (**cwnd**)**
Initial value is 1 MSS (=maximum segment size) counted as bytes
- **Actual sender window size used by TCP = minimum (**advertised win**, **cwnd**)**
- **Slow-start threshold Value (**ss_thresh**)**
Initial value is the advertised window size
- **slow start** ($\text{cwnd} < \text{ssthresh}$)
- **congestion avoidance** ($\text{cwnd} \geq \text{ssthresh}$)

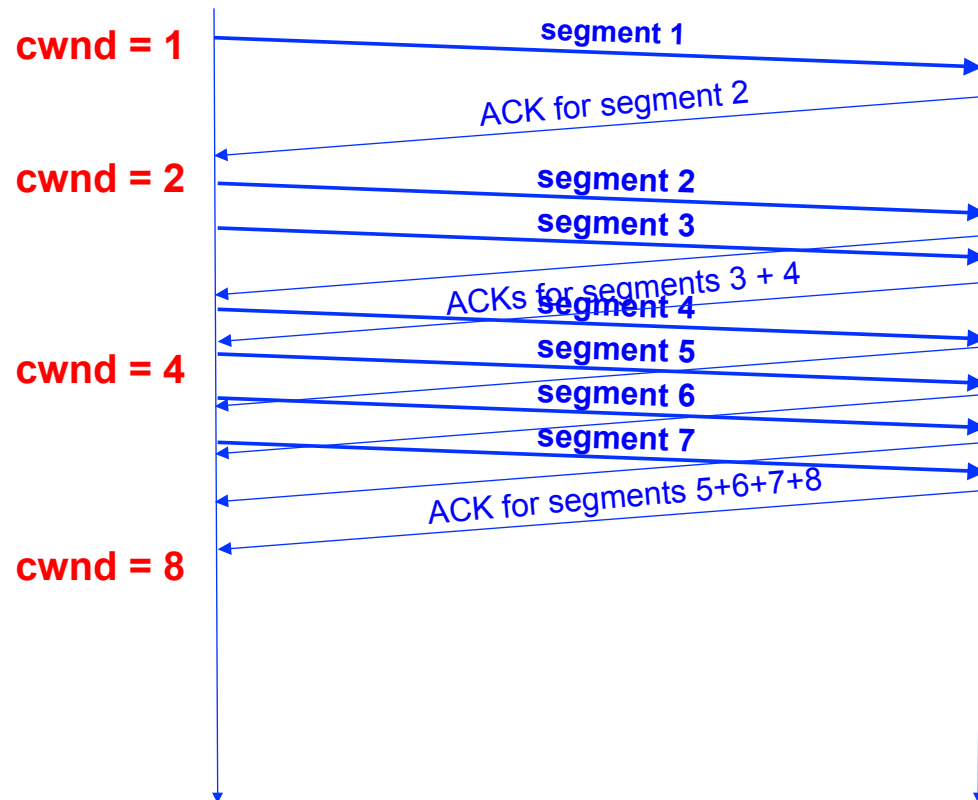
TCP: Slow Start

- Goal: discover roughly the proper sending rate quickly
- Whenever starting traffic on a new connection, or whenever increasing traffic after congestion was experienced:
 - Initialize *cwnd* = 1 MSS (max segment size)
 - Each time a segment is acknowledged, increment *cwnd* by one MSS (*cwnd*++).
- Continue until
 - Reach *ss_thresh* or
 - Packet loss



Slow Start Illustration

- The congestion window size grows very rapidly
- Observe:
 - Each ACK generates two packets
 - slow start increases rate exponentially fast (doubled every RTT)!



Congestion Avoidance

- Slow Start figures out roughly the rate at which the network starts getting congested
- Congestion Avoidance continues to react to network condition
 - Probes for more bandwidth, increase cwnd if more bandwidth available
 - If congestion detected, aggressive cut back cwnd

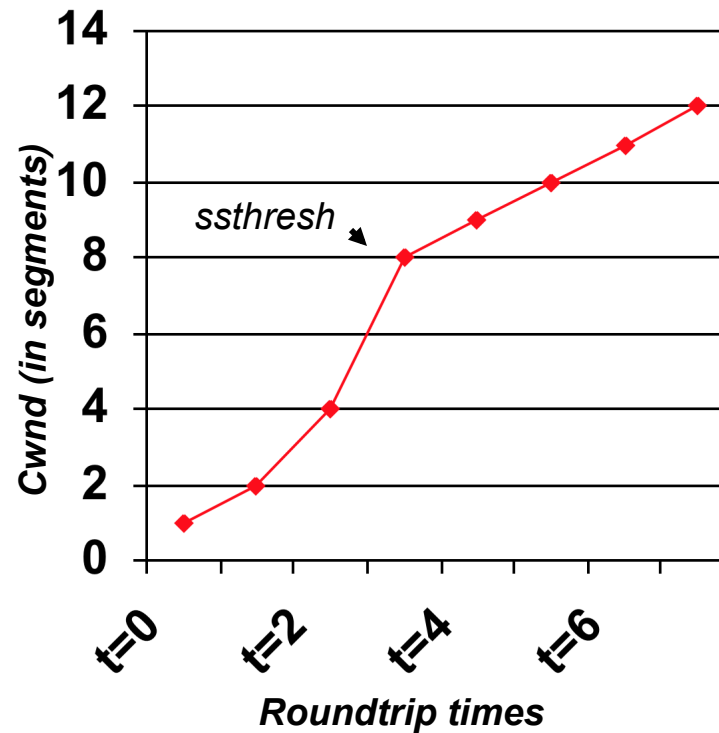


Congestion Avoidance: Additive Increase

- Slowly increase *cwnd* to probe for additional available bandwidth
- **If *cwnd* \geq *ss_thresh* then**
 each time a segment is newly acknowledged
 cwnd $+= 1/cwnd$
- *cwnd* is increased by one MSS only if all segments in the window have been acknowledged
 - Increases by 1 per RTT

Example of Slow Start + Congestion Avoidance

Assume that *ss_thresh* = 8



Detecting Congestion via Timeout


- If there is a packet loss, the ACK for that packet will not be received
- The packet will eventually timeout
 - No ack is seen as a sign of congestion

Congestion Avoidance: Multiplicative Decrease


- Each time when timeout occurs
 - ss_thresh is set to half the current size of the congestion window:
$$ss_thresh = cwnd / 2$$
 - cwnd is reset to one:
$$cwnd = 1$$
 - and slow-start is entered

Packet Pacing

- ACKs “self-pace” the data to avoid a burst of packets to be sent
- Observe: received ACK spacing \approx bottleneck bandwidth



TCP (Tahoe variant) Illustration



Many Variants of TCP

- Common variants of TCP
 - **TCP Tahoe** - the basic algorithm (discussed previously)
 - **TCP Reno** - Tahoe + fast retransmit & fast recovery
 - Many end hosts today implement TCP Reno
- and many more:
 - TCP Vegas (use timing of ACKs to avoid loss)
 - TCP SACK (selective ACK)

TCP Reno

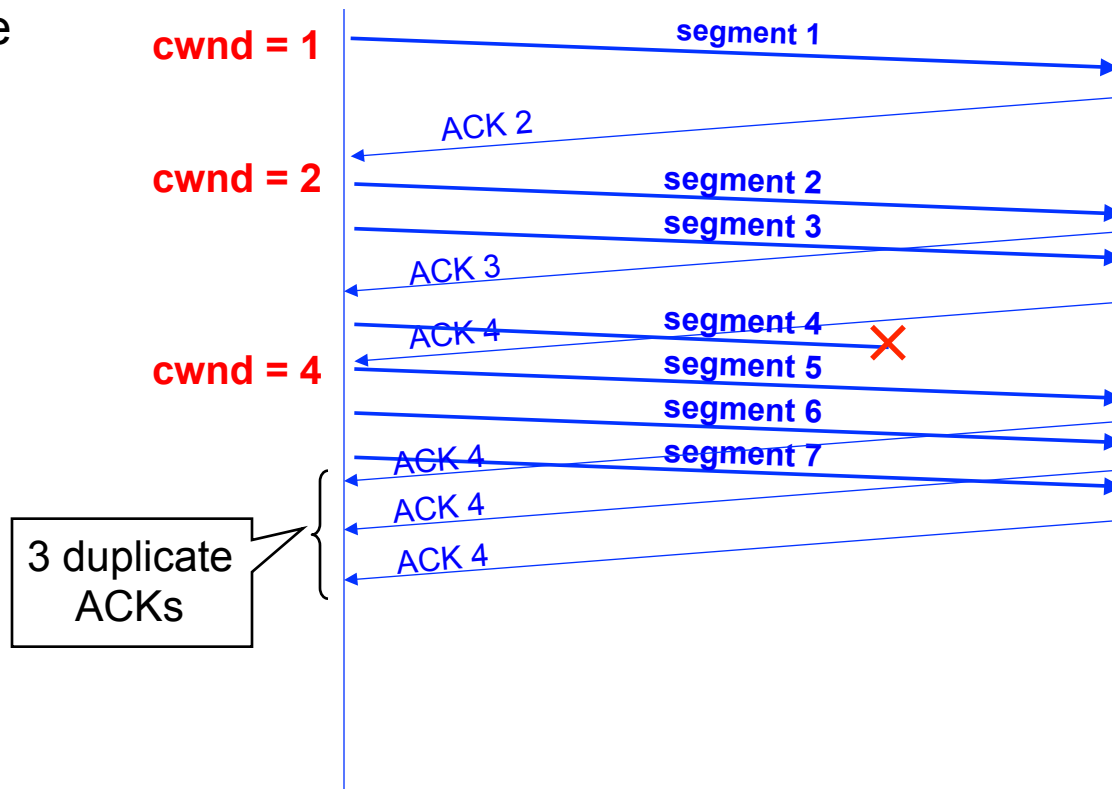
- Problem with Tahoe: If a segment is lost, there is a long wait until timeout
- Reno adds a **fast retransmit** and **fast recovery mechanisms**
- Upon receiving 3 duplicate ACKs, retransmit the presumed lost segment (“fast retransmit”)
- But do not enter slow-start. Instead enter congestion avoidance (“fast recovery”)

Fast Retransmit

- Resend a segment after 3 duplicate ACKs
 - remember a duplicate ACK means that an out-of sequence segment was received
 - ACK-n means packets 1, ..., n-1 **all** received

- Notes:

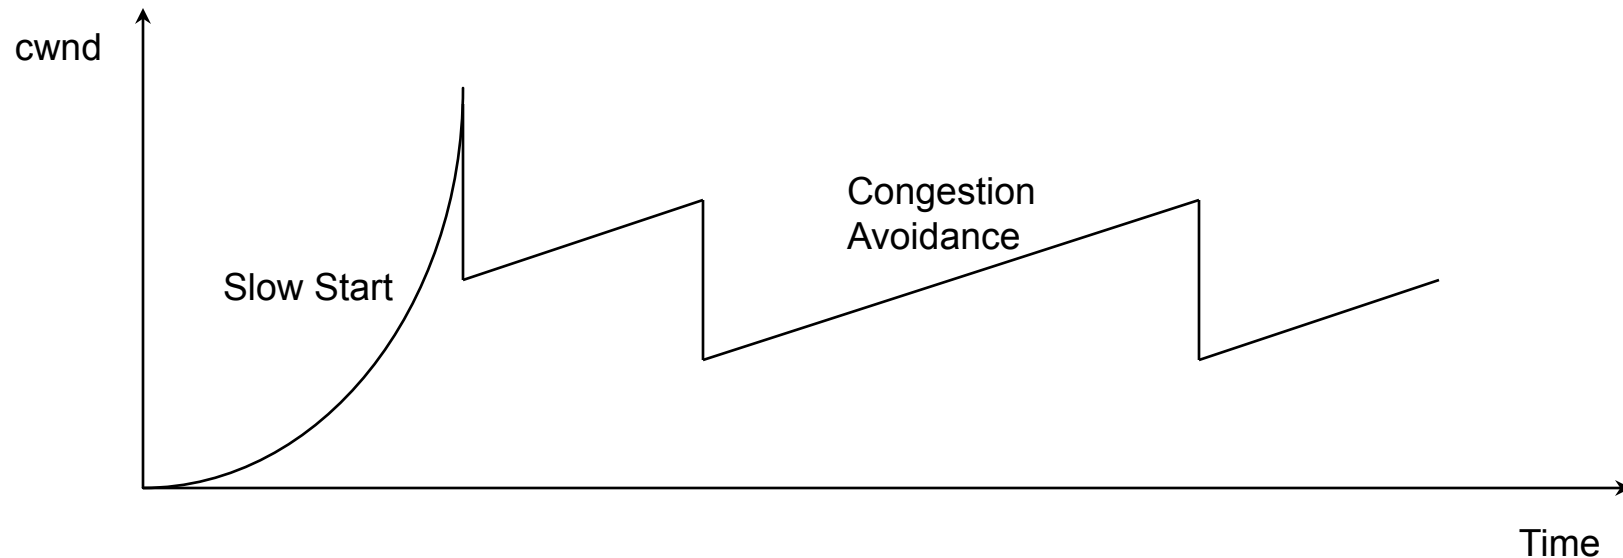
- duplicate ACKs can be due to packet reordering!
- if window is small, may not get 3 duplicate ACKs!



Fast Recovery

- After a **fast-retransmit**
 - $cwnd = cwnd/2$
 - $ss_thresh = cwnd$
 - i.e. starts congestion avoidance at new $cwnd$
- After a **timeout**
 - Same as TCP Tahoe
 - $ss_thresh = cwnd/2$
 - $cwnd = 1$
 - Do slow start

Fast Retransmit and Fast Recovery



- Retransmit after 3 duplicate ACKs
 - prevent expensive timeouts
- Slow start only once per session (if no timeouts)
- In steady state, *cwnd* oscillates around the ideal window size.

TCP Reno Summary

- Slow-Start if $\text{cwnd} < \text{ss_thresh}$
 - $\text{cwnd}++$ upon every new ACK (exponential growth)
 - Timeout: $\text{ss_thresh} = \text{cwnd}/2$ and $\text{cwnd} = 1$
- Congestion avoidance if $\text{cwnd} \geq \text{ss_thresh}$
 - Additive Increase Multiplicative Decrease (AIMD)
 - ACK: $\text{cwnd} = \text{cwnd} + 1/\text{cwnd}$
 - Timeout: $\text{ss_thresh} = \text{cwnd}/2$ and $\text{cwnd} = 1$
- Fast Retransmit & Recovery
 - 3 duplicate ACKs (interpret as packet loss)
 - Retransmit lost packet
 - $\text{cwnd} = \text{cwnd}/2$, $\text{ss_thresh} = \text{cwnd}$

TCP Reno Saw Tooth Behavior

