

# COMP/ELEC 429/556

## Introduction to Computer Networks

Network security

Some slides used with permissions from Edward W.  
Knightly, T. S. Eugene Ng, Ion Stoica, Hui Zhang

## The 2004 A. M. Turing Award Goes to...



Bob Kahn

Vint Cerf

- *"For pioneering work on internetworking, including the design and implementation of the Internet's basic communications protocols, TCP/IP, and for inspired leadership in networking."*

# Design Philosophy of the DARPA Internet Protocols by David D. Clark (1988)

1. Internet communication must continue despite loss of networks or gateways
2. The Internet must provide a wide range of communications services
3. The Internet must support a wide variety of applications
4. The Internet must support a wide range of management functions
5. The Internet must support a wide range of services
6. The Internet must support a wide range of services with a low level of overhead
7. The resources must be accounted for



**FRAGILE**

# Importance of Network Security

- Internet currently used for important services
  - Financial transactions, medical records
- Could be used in the future for *critical* services
  - 911, surgical operations, energy system control, transportation system control
- Networks more open than ever before
  - Global, ubiquitous Internet; wireless access
- Malicious Users
  - Selfish users: want more network resources
  - Malicious users: want to do harm

# Network Security: Our Focus

- Host Compromise
  - Attacker gains control of a host
- Denial-of-Service
  - Attacker prevents legitimate users from gaining service
- Note: Attack can be both
  - E.g., host compromise that provides resources for denial-of-service

# Host Compromise Method 1: User Exploitation

- Many applications rely on the user to decide if a potentially dangerous action should be taken, e.g.,
  - Run code downloaded from the Internet
    - “Do you accept content from Microsoft?”
  - Run code attached to email
    - “Subject: You’ve got to see this!”
  - Allow a macro in a data file to be run
    - “Here is the latest version of the document.”

## Host Compromise Method 2: Stack Based Buffer Overflow (User not involved!)


- Code can have many unknown bugs because those bugs are not triggered by common input
- Such bugs in network facing code can be triggered by inputs received from the network
- Network facing code that runs with high privileges (i.e., as root) is especially dangerous

## Example

- What is wrong here?

```
#define MAXNAMELEN 64
int offset = 4;
char username[MAXNAMELEN];
int name_len;

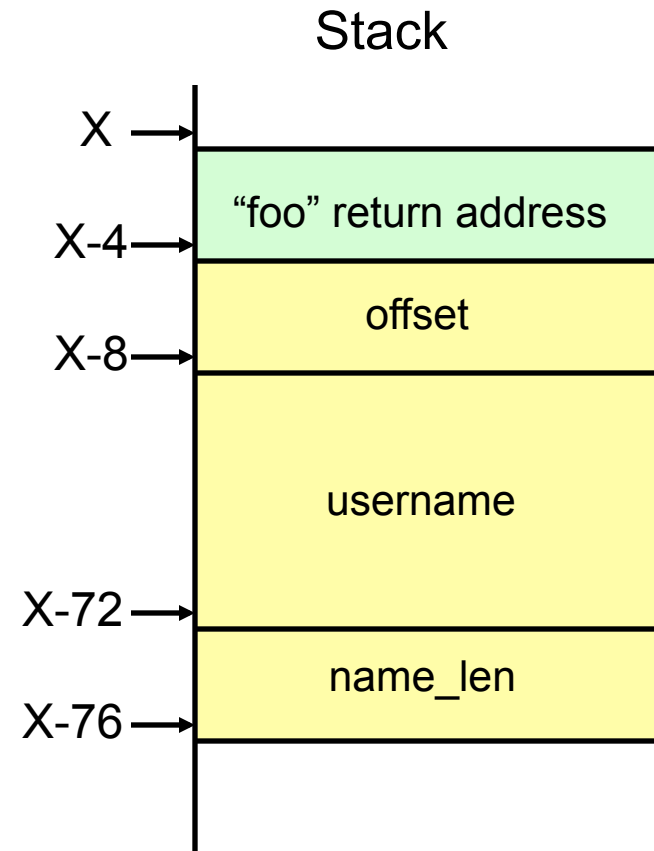
name_len = ntohl(*(int *)packet);
memcpy(&username, packet[offset], name_len);
```





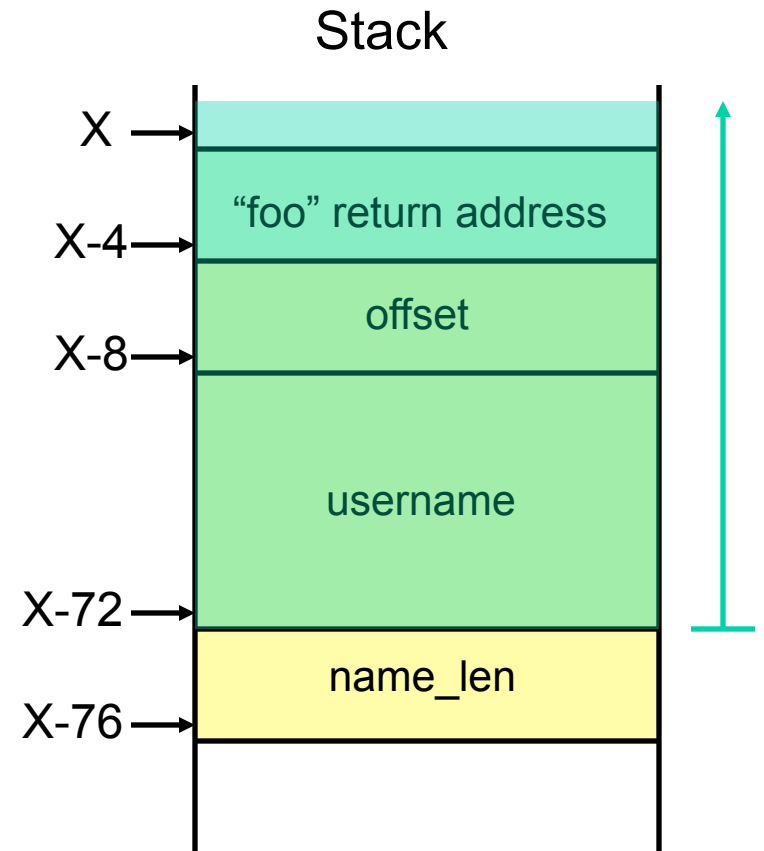
# Example

```
void foo(packet) {  
    #define MAXNAMELEN 64  
    int offset = 4;  
    char username[MAXNAMELEN];  
    int name_len;  
  
    name_len = ntohl(*(int*)packet);  
    → memcpy(&username,  
            packet[offset], name_len);  
    ...  
}
```



# Example

```
void foo(packet) {  
    #define MAXNAMELEN 64  
    int offset = 4;  
    char username[MAXNAMELEN];  
    int name_len;  
  
    name_len = ntohl(*(int *) packet);  
    memcpy(&username,  
           packet[offset], name_len);  
    ...  
}
```



# Hall of Shame

- Software that have had many stack overflow bugs:
  - BIND (a popular DNS server)
  - RPC (Remote Procedure Call, used for Network File System)
  - Sendmail (a popular UNIX mail delivery software)
  - IIS (Windows web server)
  - SNMP (Simple Network Management Protocol, used to manage routers and other network devices)

# Effect of Stack Based Buffer Overflow

- Write into part of the stack or heap
  - Write arbitrary code to part of memory
  - Cause program execution to jump to arbitrary code
- Worm
  - Sends bogus input to potential victims to spread the worm
  - Attacker can do anything that the privileges of the buggy program allows
    - Launches copy of itself on compromised host



# Internet Worm

- One of earliest major Internet security incidents
  - Internet Worm (1988): compromised almost every BSD-derived machine on Internet
- Today: estimated that a single worm could compromise 10M hosts in  $< 5$  min
- Attacker gains control of a host
  - Reads data
  - Erases data
  - Compromises another host
  - Launches denial-of-service attack on another host

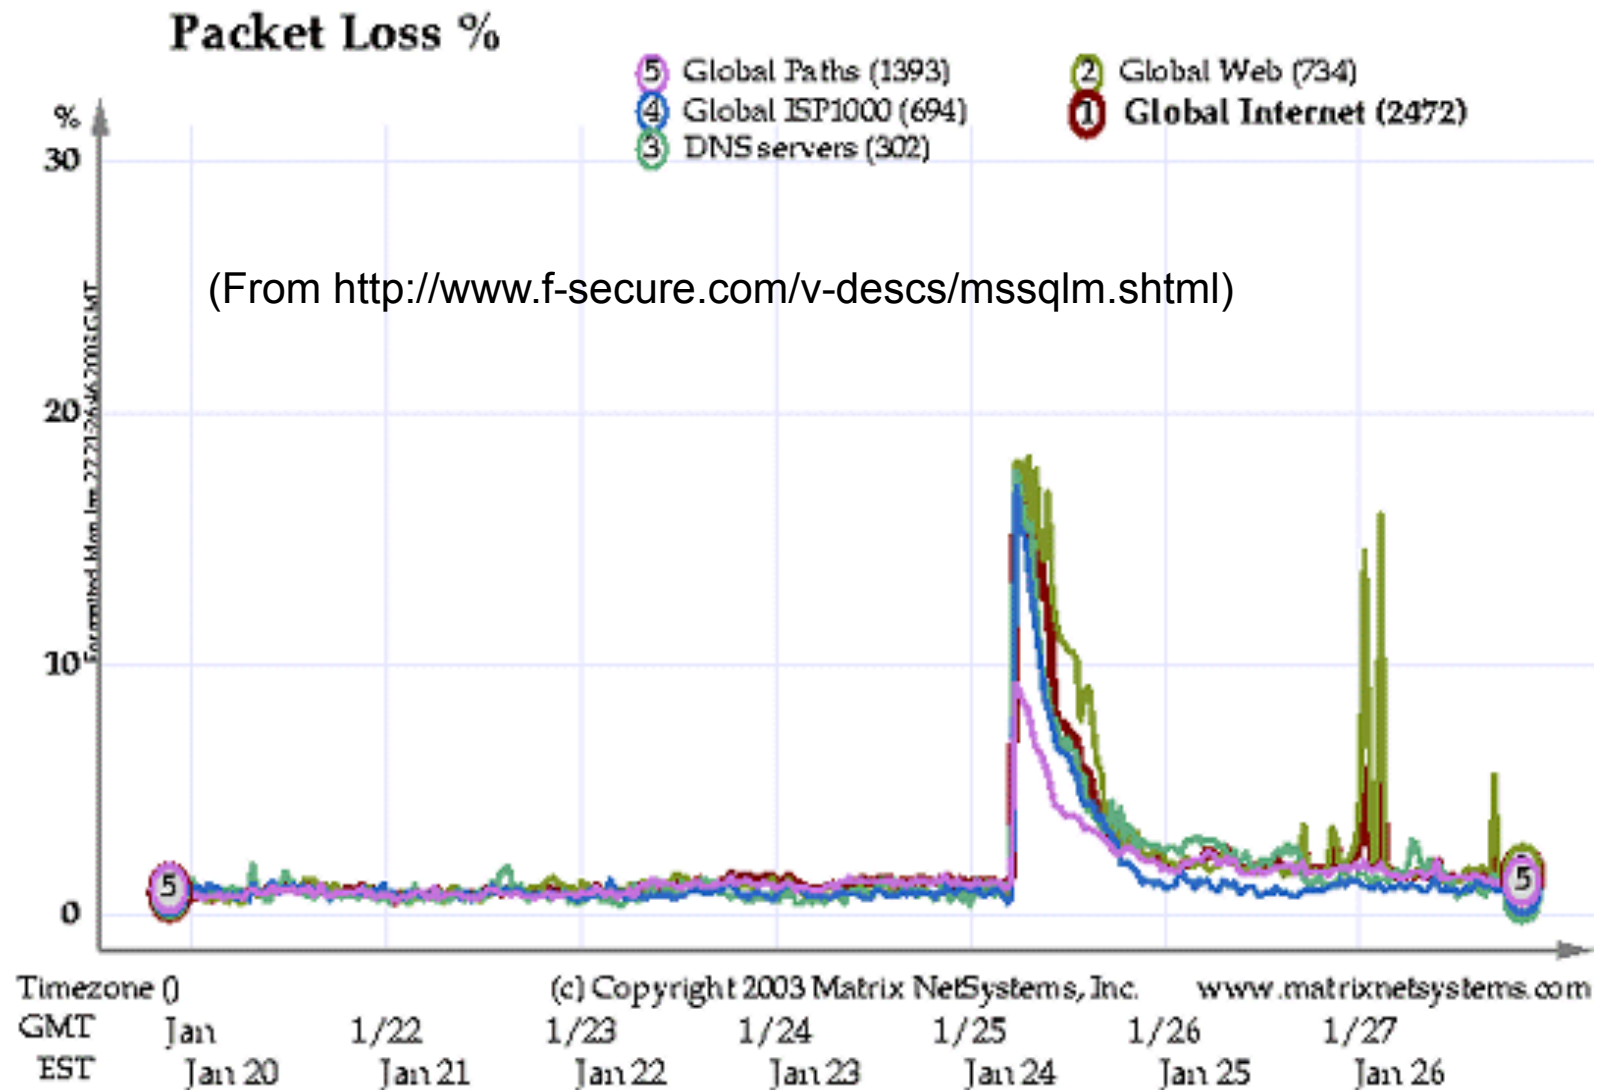
# Worm Spreading

$$f = (e^{K(t-T)} - 1) / (1 + e^{K(t-T)})$$

- $f$  – fraction of hosts infected
- $K$  – rate at which one host can contact others
- $T$  – start time of the attack



## e.g. MS SQL Slammer (January 2003)



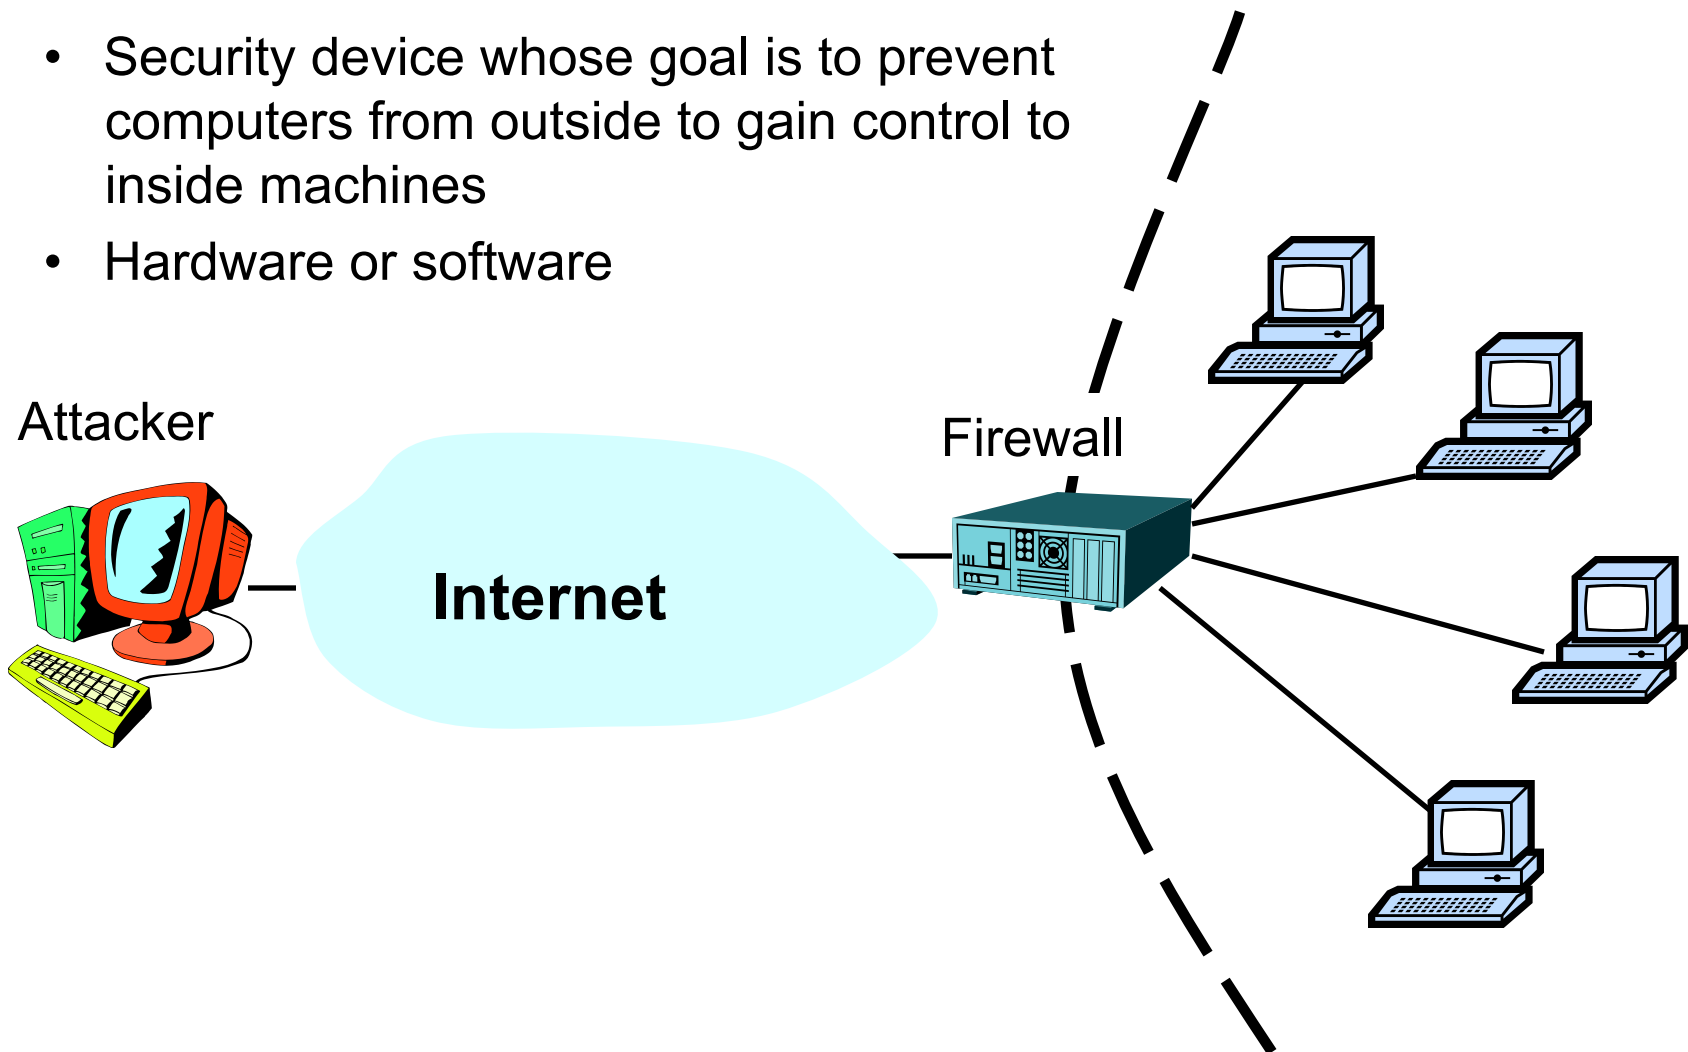
## Potential Solutions

- Don't write buggy software
- Type-safe Languages
  - Unrestricted memory access of C/C++ contributes to problem
  - Use Java, Python, etc. instead (but there is a performance cost)
- Operating system
  - Make stack memory pages non-executable
  - Add a guard next to the return address on stack, check guard value before executing jump instruction
  - Compartmentalize programs better, so one compromise doesn't compromise the entire system
    - E.g., DNS server doesn't need total system access
- Firewalls



# Firewall

- Security device whose goal is to prevent computers from outside to gain control to inside machines
- Hardware or software



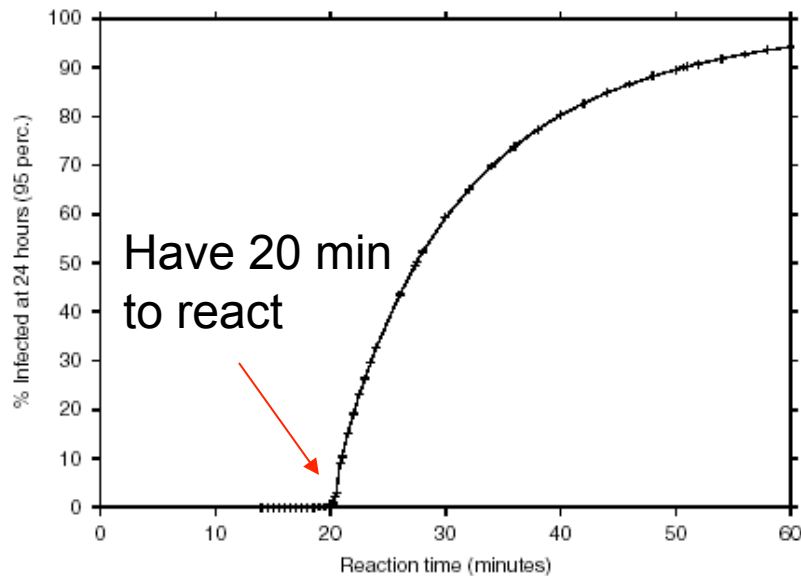
# Firewall

- Restrict traffic between Internet and devices (machines) behind it based on
  - Source address and port number
  - Payload
  - Stateful analysis of data
- Examples of rules
  - Block any external packets not for port 80
  - Block any email with an attachment
  - Block any external packets with an internal IP address as source


## Firewall Properties

- Easier to deploy firewall than secure all internal hosts
- Doesn't prevent user exploitation
- Can't prevent problem from spreading from within
- Tradeoff between availability of services (firewall allows/disallows ports) and security
  - If firewall is too restrictive, users will find way around it, thus compromising security
  - E.g., have all services use port 80, which is typically not blocked by firewalls

# Address Blacklisting and Content Filtering at Firewalls: Code Red Worm Example



**Address Blacklisting:** Set firewalls to block an IP address after it has been detected as malicious



**Content Filtering:** Set firewalls to block packets with malicious contents after such contents are identified

Bottom line: Internet worms are very hard  
to contain

Compromised hosts are then used to  
launch Denial of Service attacks

# Denial of Service

- Huge problem in current Internet
  - Major sites attacked: Google, Yahoo!, Amazon, eBay, CNN, Microsoft, ...
  - Almost all attacks launched from compromised hosts
    - So called Botnets
- Impact
  - Prevent legitimate users from gaining service by overloading or crashing the service

## Exploit “Asymmetry”

- Attacker possesses asymmetric amount of power to disrupt victim than the victim’s power to defend itself
- E.g. Use a large number of compromised hosts
- E.g. Exploit “asymmetry” in protocol designs and in software implementations
  - Focus of our discussion


## E.g. Ping Flood

- Attacker sends ICMP PING packets with IP subnet broadcast address as the destination address
  - e.g. 128.42.255.255
- Last hop router broadcasts the PING packet in the victim's Ethernet network
- Amplify attacker's power to congest the victim's Ethernet network




## E.g. Reflection

- Reflection
  - Cause one non-compromised host to attack another
  - E.g., host A sends DNS request with source V to server R. R sends reply to V



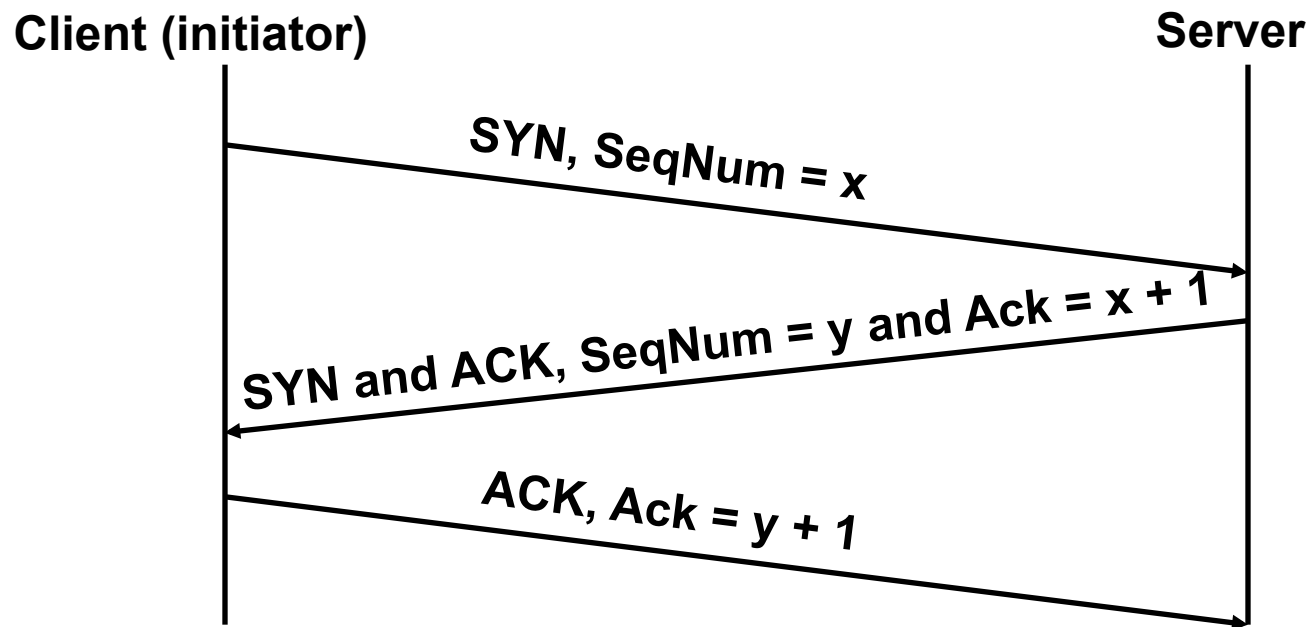
## E.g. Reflection

- Reflection
  - Cause one non-compromised host to attack another
  - E.g., host A sends DNS request with source V to server R. R sends reply to V



## E.g. SYN Attack (Recap: 3-Way Handshaking)

- Goal: agree on a set of parameters: the start sequence number for each side
  - Starting sequence numbers are random.



## E.g. SYN Attack

- Attacker: send at max rate TCP SYN with random spoofed source address to victim
  - Spoofing: use a different source IP address than own
  - Random spoofing allows one host to pretend to be many
- Victim receives many SYN packets
  - Send SYN+ACK back to spoofed IP addresses
  - Holds some memory until 3-way handshake completes
    - Usually never, so victim times out after long period (e.g., 3 minutes)


## Effect on Victim

- Buggy implementations allow unfinished connections to eat all memory, leading to crash
- Better implementations limit the number of unfinished connections
  - Once limit reached, new SYNs are dropped
- Effect on victim's users
  - Users can't access the targeted service on the victim because the unfinished connection queue is full → DoS

## Solution: SYN Cookies

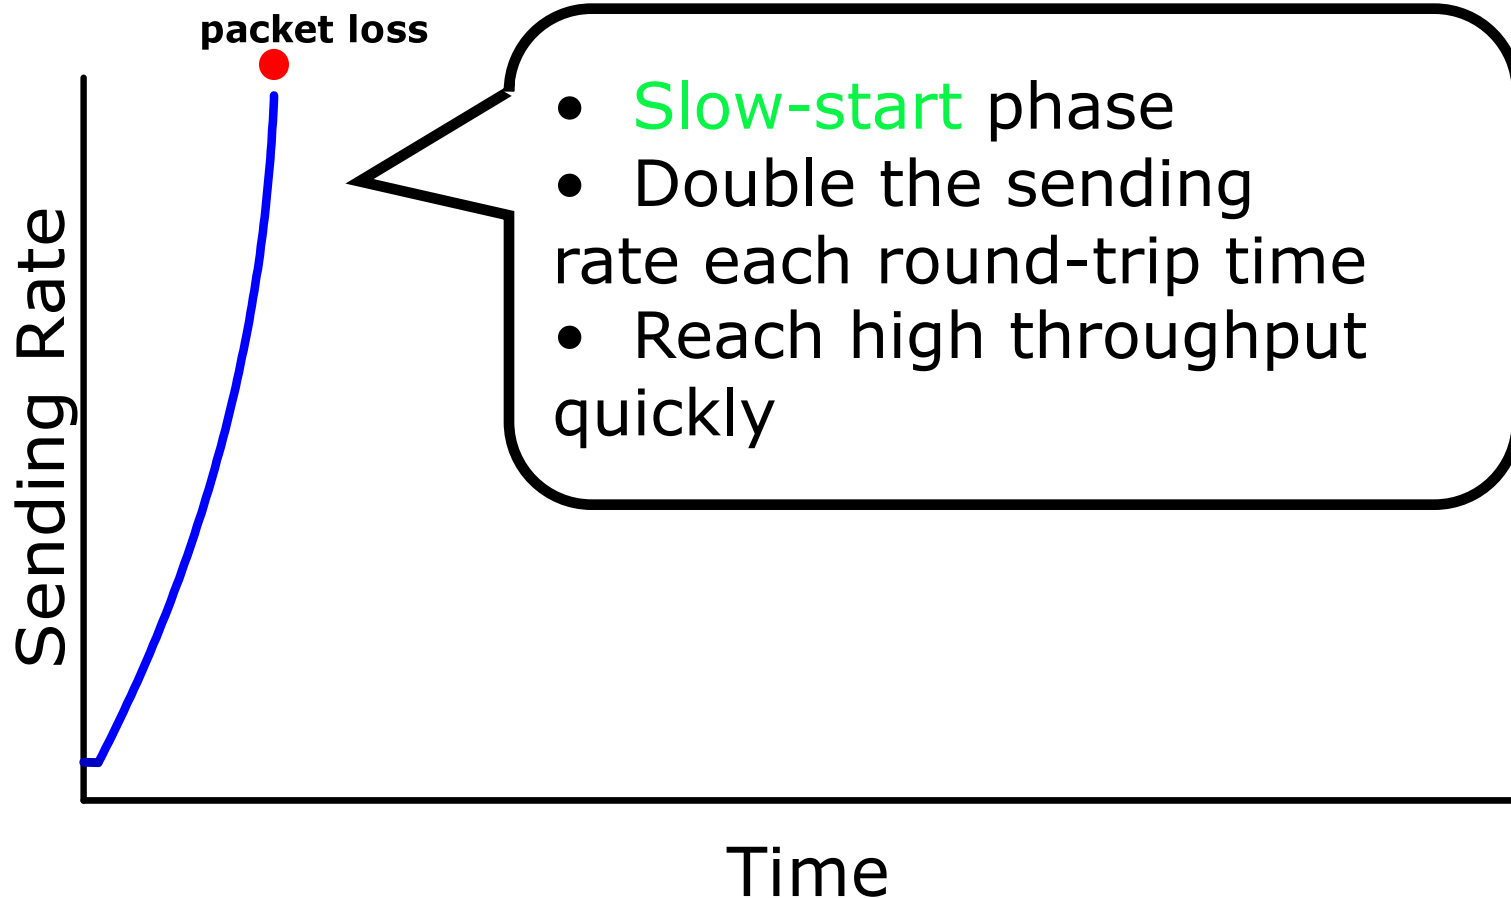
- Server: send SYN-ACK with sequence number  $y$ , where
  - $y = H(\text{client\_IP\_addr}, \text{client\_port}, \text{server\_secret})$
  - $H()$ : one-way hash function
- Client: send ACK containing  $y+1$
- Server:
  - verify if  $y = H(\text{client\_IP\_addr}, \text{client\_port}, \text{server\_secret})$
  - If verification passes, allocate memory
- Note: server doesn't allocate any memory if the client's address is spoofed

E.g. “Shrew attack” (Low rate TCP attack)  
by former ECE student Aleksander  
Kuzmanovic



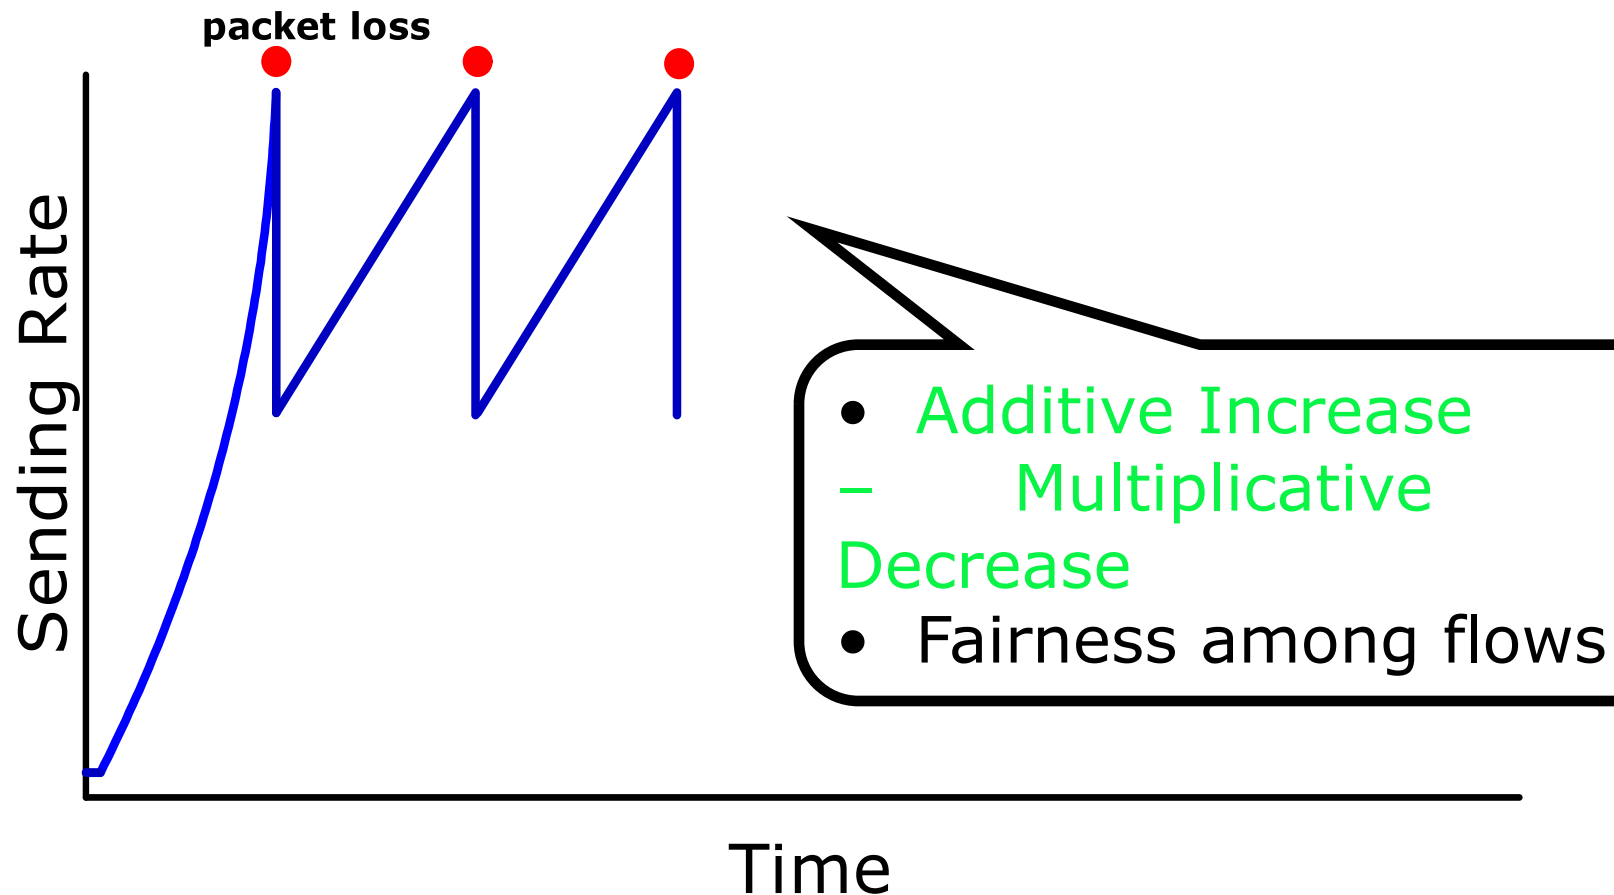
- Very small but aggressive mammal that ferociously attacks and kills much larger animals with a venomous bite

# TCP Congestion Control

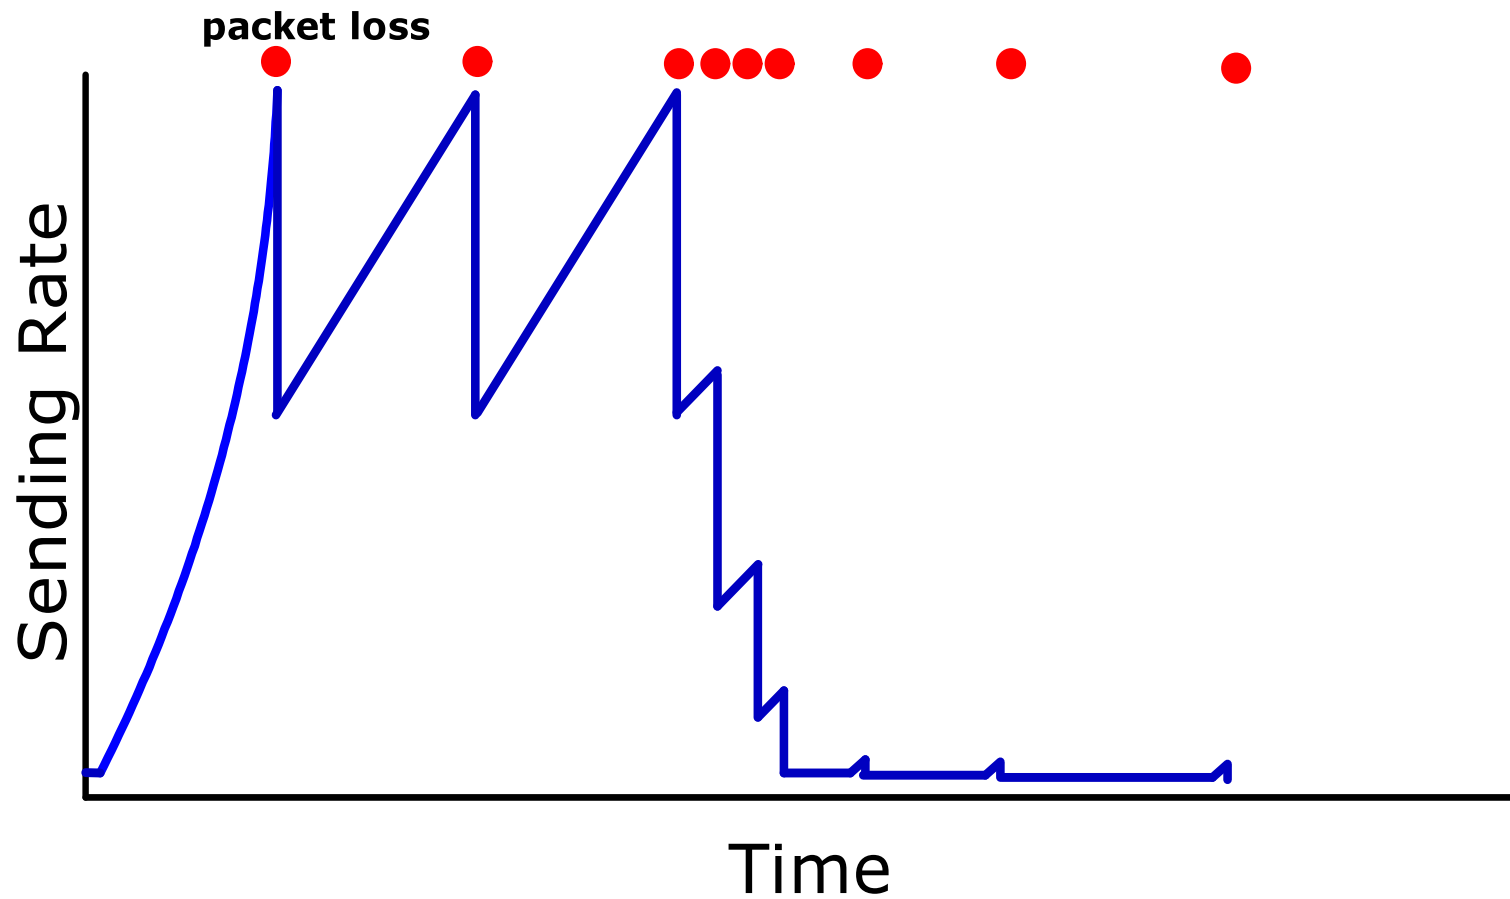




# TCP Congestion Control




# TCP Congestion Control






# TCP Retransmission Timeout (RTO)

- $RTO = SRTT + 4 * RTTVAR$ 
  - RTT ~ 10-100ms
- Lower-bounding the RTO parameter:
  - [AllPax99]: minRTO = 1 sec
    - to avoid spurious retransmissions
  - RFC6298 recommends minRTO = 1 sec
    - Many implementations actually use 200ms

# The Shrew Attack






# The Shrew Attack






- A short burst ( $\sim$ RTT) sufficient to create outage
- **Outage** – event of correlated packet losses that forces TCP to enter RTO mechanism

# The Shrew Attack





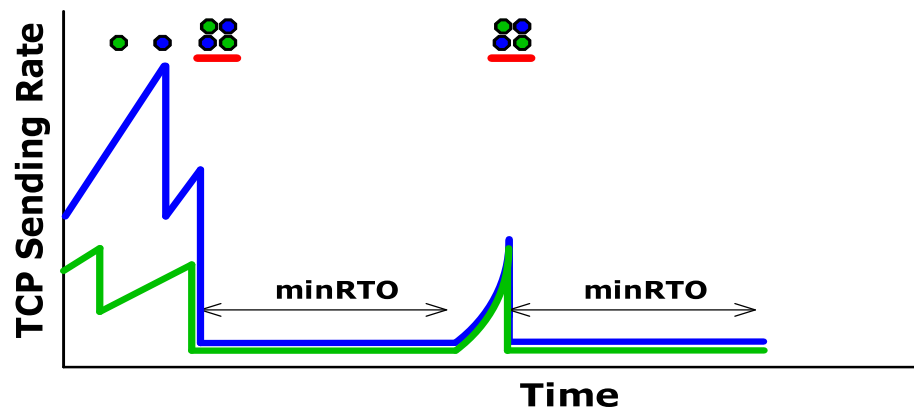
- The outage **synchronizes** all TCP flows
  - All flows react **simultaneously** and **identically**
    - backoff for minRTO

# The Shrew Attack



- Once the TCP flows try to recover – hit them again
- Exploit protocol **determinism**

# The Shrew Attack



- And keep repeating...
- RTT-time-scale outages interspaced on minRTO periods can deny service to TCP traffic



## Shrew Principles

- A single RTT-length outage forces all TCP flows to *simultaneously* enter timeout
  - All flows respond *identically* and stop for the minRTO period
  - Most invoke slow start afterwards
- Shrews exploit protocol *determinism*, and repeat the outage after each minRTO period
- Periodic outages *synchronize* TCP flows and deny their service
- Outages occur relatively slowly and can be induced with low average rate

# Difficulties of dealing with DoS Attacks

- Distinguish DoS attack from flash crowd
  - Can be very hard if the attack is sophisticated
- Prevent damage
  - Distinguish (if possible) attack traffic from legitimate traffic
  - Rate limit attack traffic
- Stop attack
  - Identify attacking machines, not easy if source address is spoofed
  - Shutdown attacking machines
  - Usually done manually, requires cooperation of ISPs
- Identify persons responsible
  - Very difficult
  - Usually done manually, requires cooperation of ISPs

## Prevent Damage

- Use fair queuing to limit damage if attack traffic is not distinguishable
- Prevent an attacker from sending at 10Mb/s and hurting a user sending at 1Mb/s
- Does not prevent 10 attackers from sending at 1Mb/s and hurting a user sending a 1Mb/s

## Stop Attack

- Need to identify attacking machines, but not easy if source addresses are spoofed

Some ideas to defeat spoofed source addresses:

- Ingress filtering
  - A domain's border router drop outgoing packets which do not have a valid source address for that domain
  - If universally used, could abolish spoofing
- IP Traceback
  - Routers probabilistically tag packets with an identifier
  - Destination can infer path to true source after receiving enough packets

## In Closing...

- You can now explain to your Mom and Dad how the Internet works!!
- Application level
  - Domain name/IP address conversion, socket programming (TCP and UDP, client/server), buffer overflow bug, worm, DoS attack
- Transport level
  - Reliability (sliding window, CRC, etc.), congestion control (AIMD)
- Network level
  - Hierarchical addressing, intra-domain routing (LS and DV), inter-domain routing (path vector, policies)
- Link level
  - Broadcast network access control (Aloha, CSMA/CD, WiFi), Ethernet spanning tree, weighted fair queuing
- Physical level
  - Bit encoding (NRZI, 4B/5B, etc), packet framing (bit and byte stuffing)